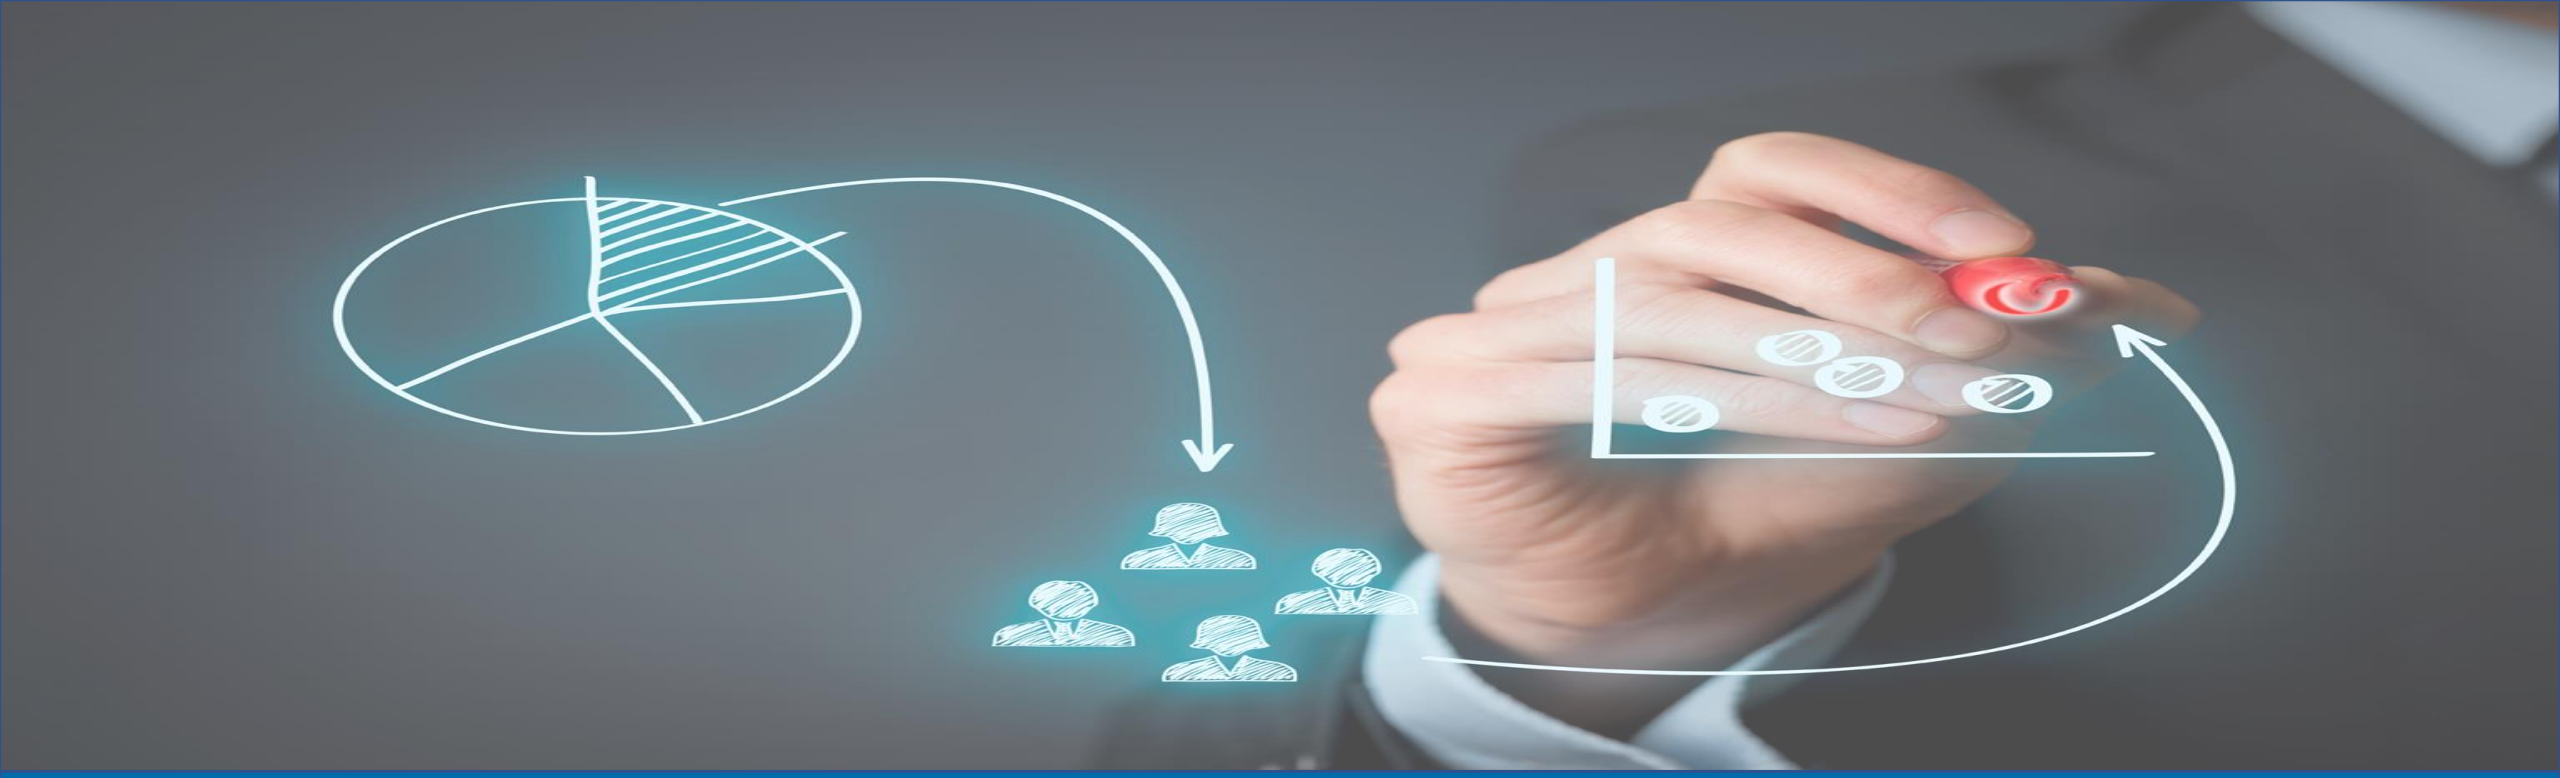

PROCESSING BIG DATA

with Azure Data Lake Analytics

Sean Forgatch
Business Intelligence Consultant
Sean.Forgatch@talavant.com

TALAVANT

About Me

Sean Forgatch

- Milwaukee, WI
- Business Intelligence Consultant
  - Healthcare, Insurance, SaaS
  - Integration and Analytics
  - Microsoft Big Data Certified
- PASS
  - Industry Speaker
  - FoxPASS President
- Running, Craft Beers, Reading

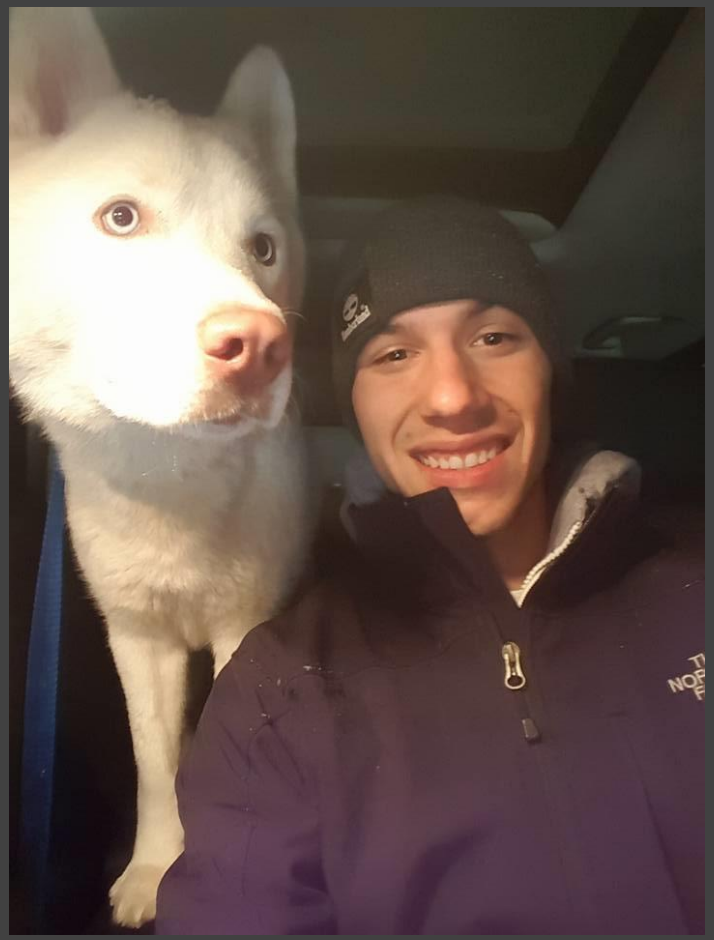

About Talavant

There is a better way to make data work for companies. Better resources, strategy, sustainability, inclusion of the organization as a whole, understanding of client needs, tools, outcomes, better ROI.

STRATEGY

ARCHITECTURE

IMPLEMENTATION

VALUE WE PROVIDE

- Accelerated planning, implementation and results
- Sustainable
- Increased

HOW WE DO IT

By providing a holistic approach inclusive of a client's people, processes and technologies - built on investment in our own employees and company growth.

TALAVANT

1. **Big Data Overview**
2. Data Lake Concepts
3. Azure Data Lake Store
4. Azure Data Lake Analytics
5. U-SQL

TALAVANT

GDELT Analysis in PowerBI

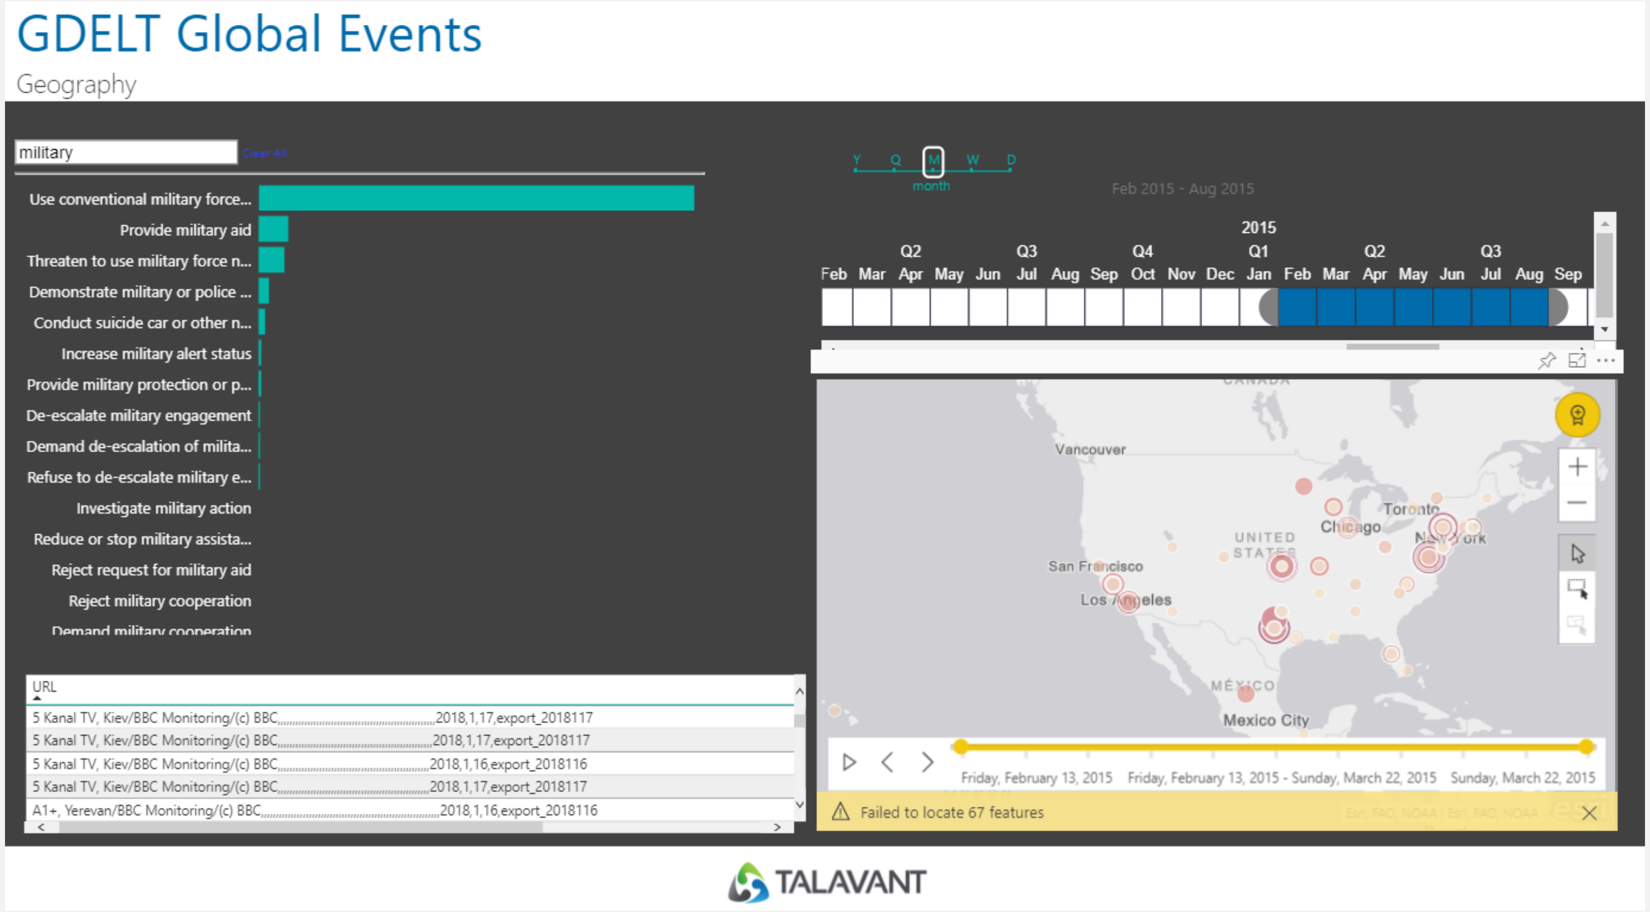

Big Data Primer

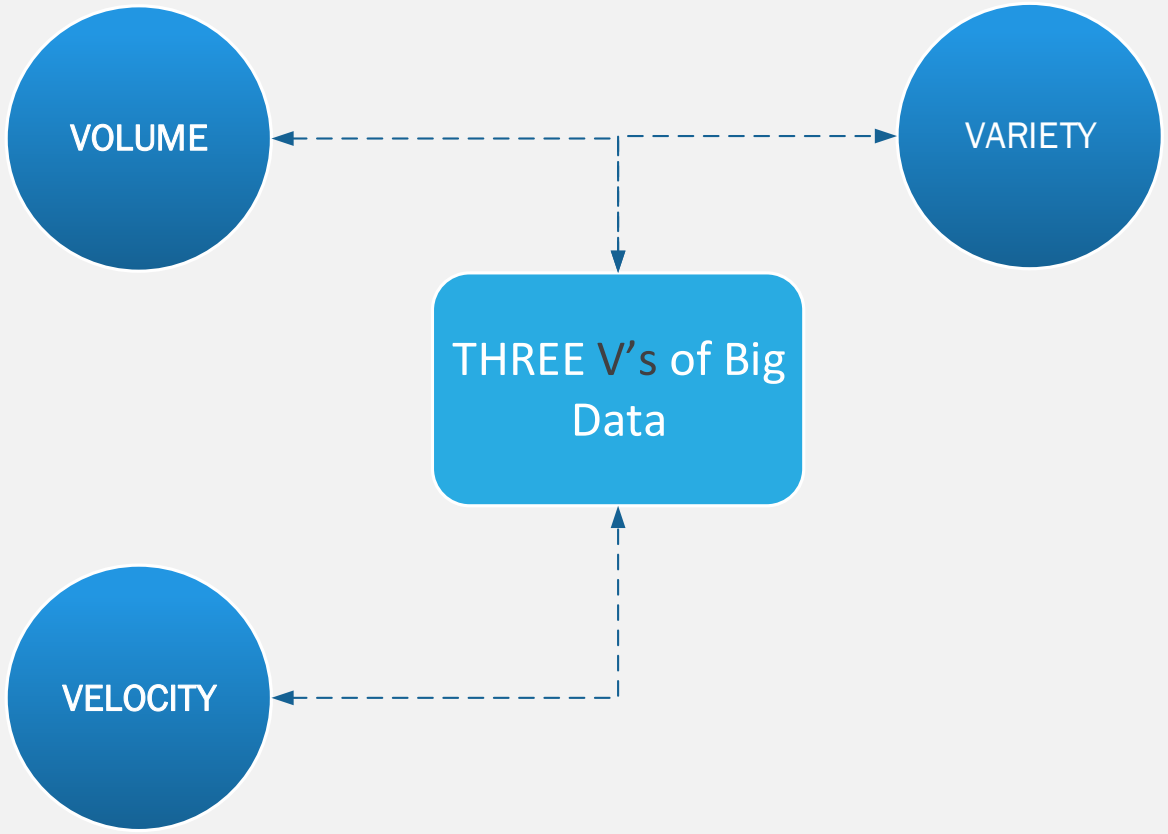

Big Data Primer: Azure Tools

A Big Data Trends

“Variety, not volume or velocity, drives big-data investments”

A Big Data Trends

Variety, not volume or velocity, drives big-data investments”

“Big data grows up: Hadoop adds to enterprise standards”

A Big Data Trends

“Variety, not volume or velocity, drives big-data investments”

“Big data grows up: Hadoop adds to enterprise standards”

“Rise of metadata catalogs helps people find analysis-worthy big data.”

-TDWI: Top Ten Big Data Trends for 2017

1. Big Data Overview
2. **Data Lake Concepts**
3. Azure Data Lake Store
4. Azure Data Lake Analytics
5. U-SQL

Data Lake Concepts

Data Lake Operations

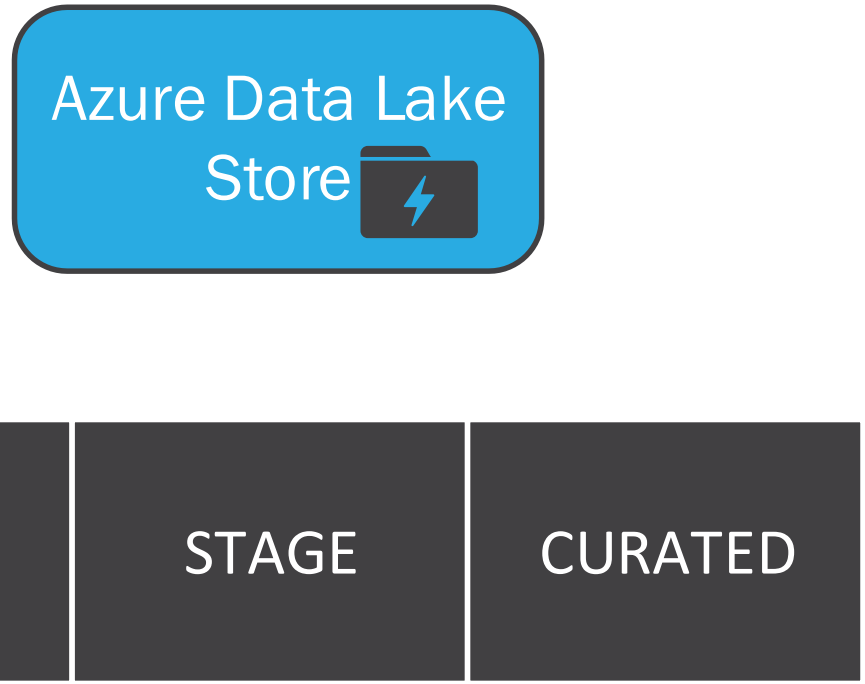

Data Lake Operations

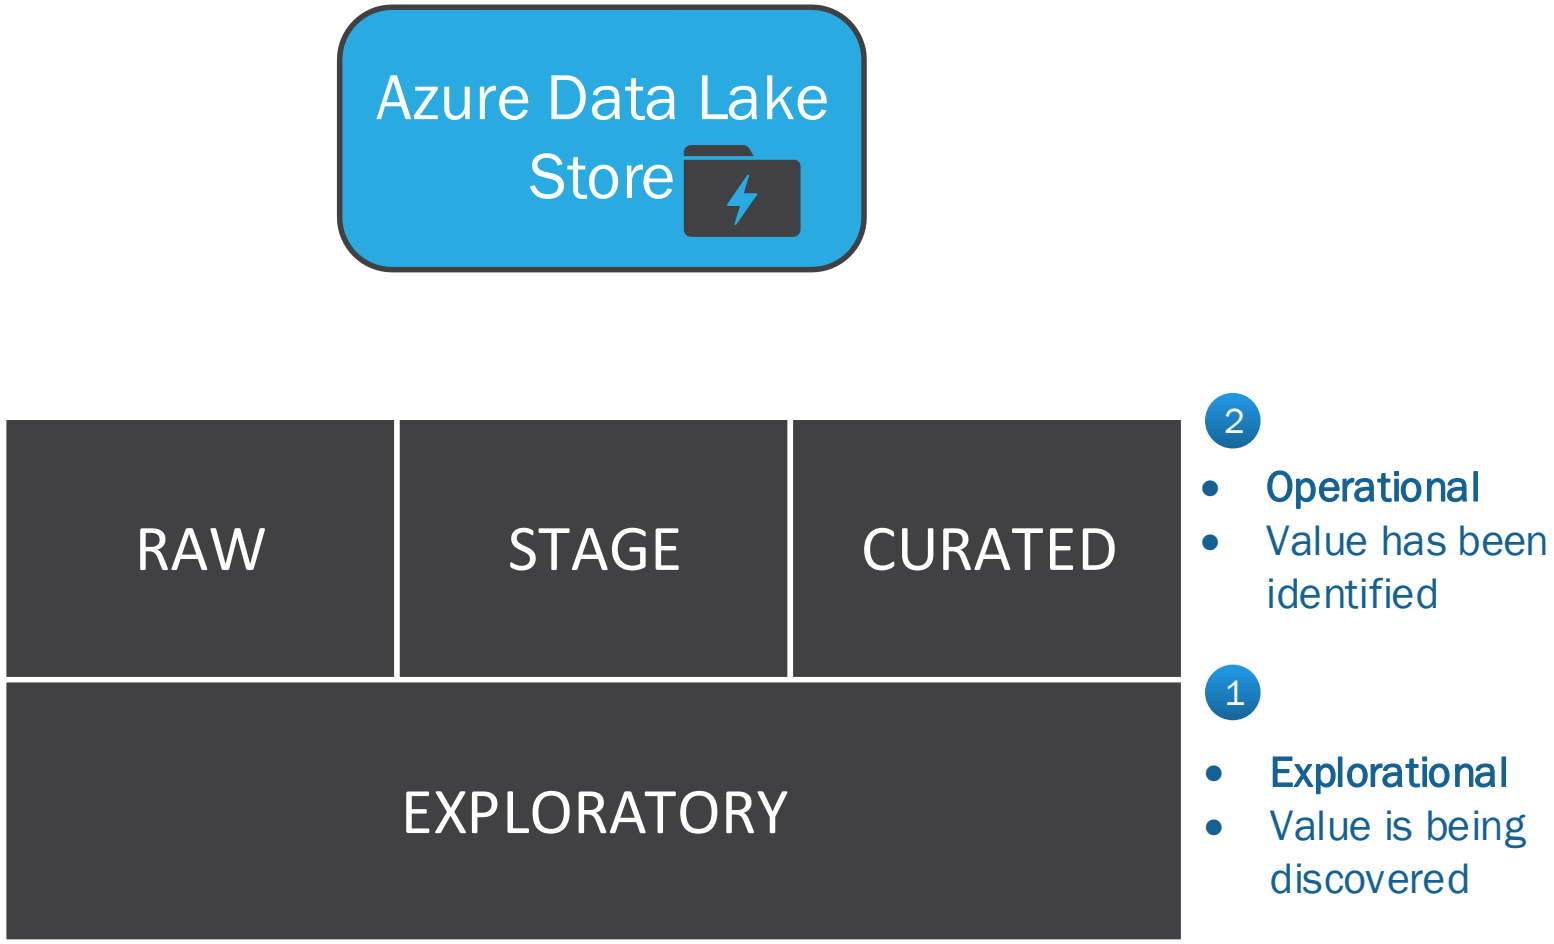

Data Lake Operations

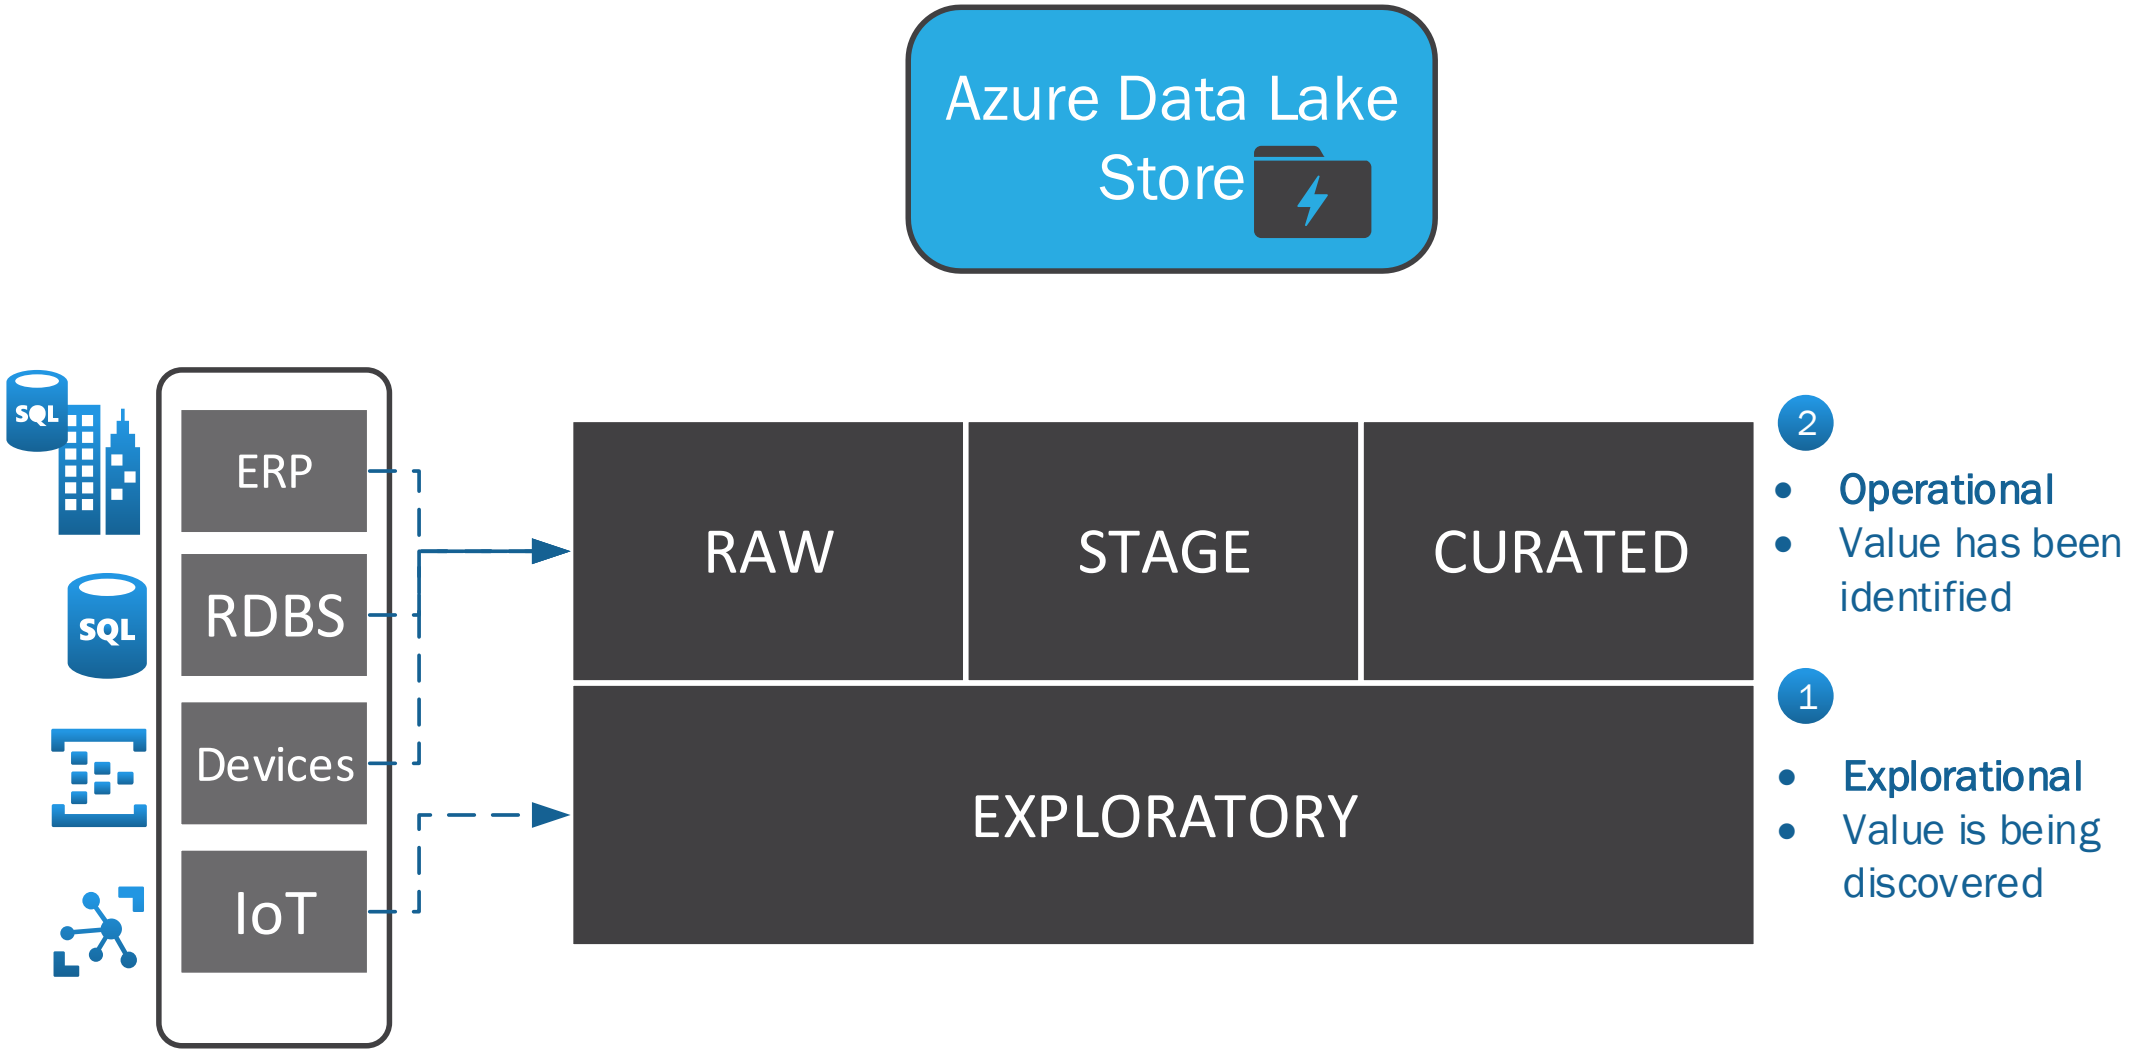

Data Lake Operations

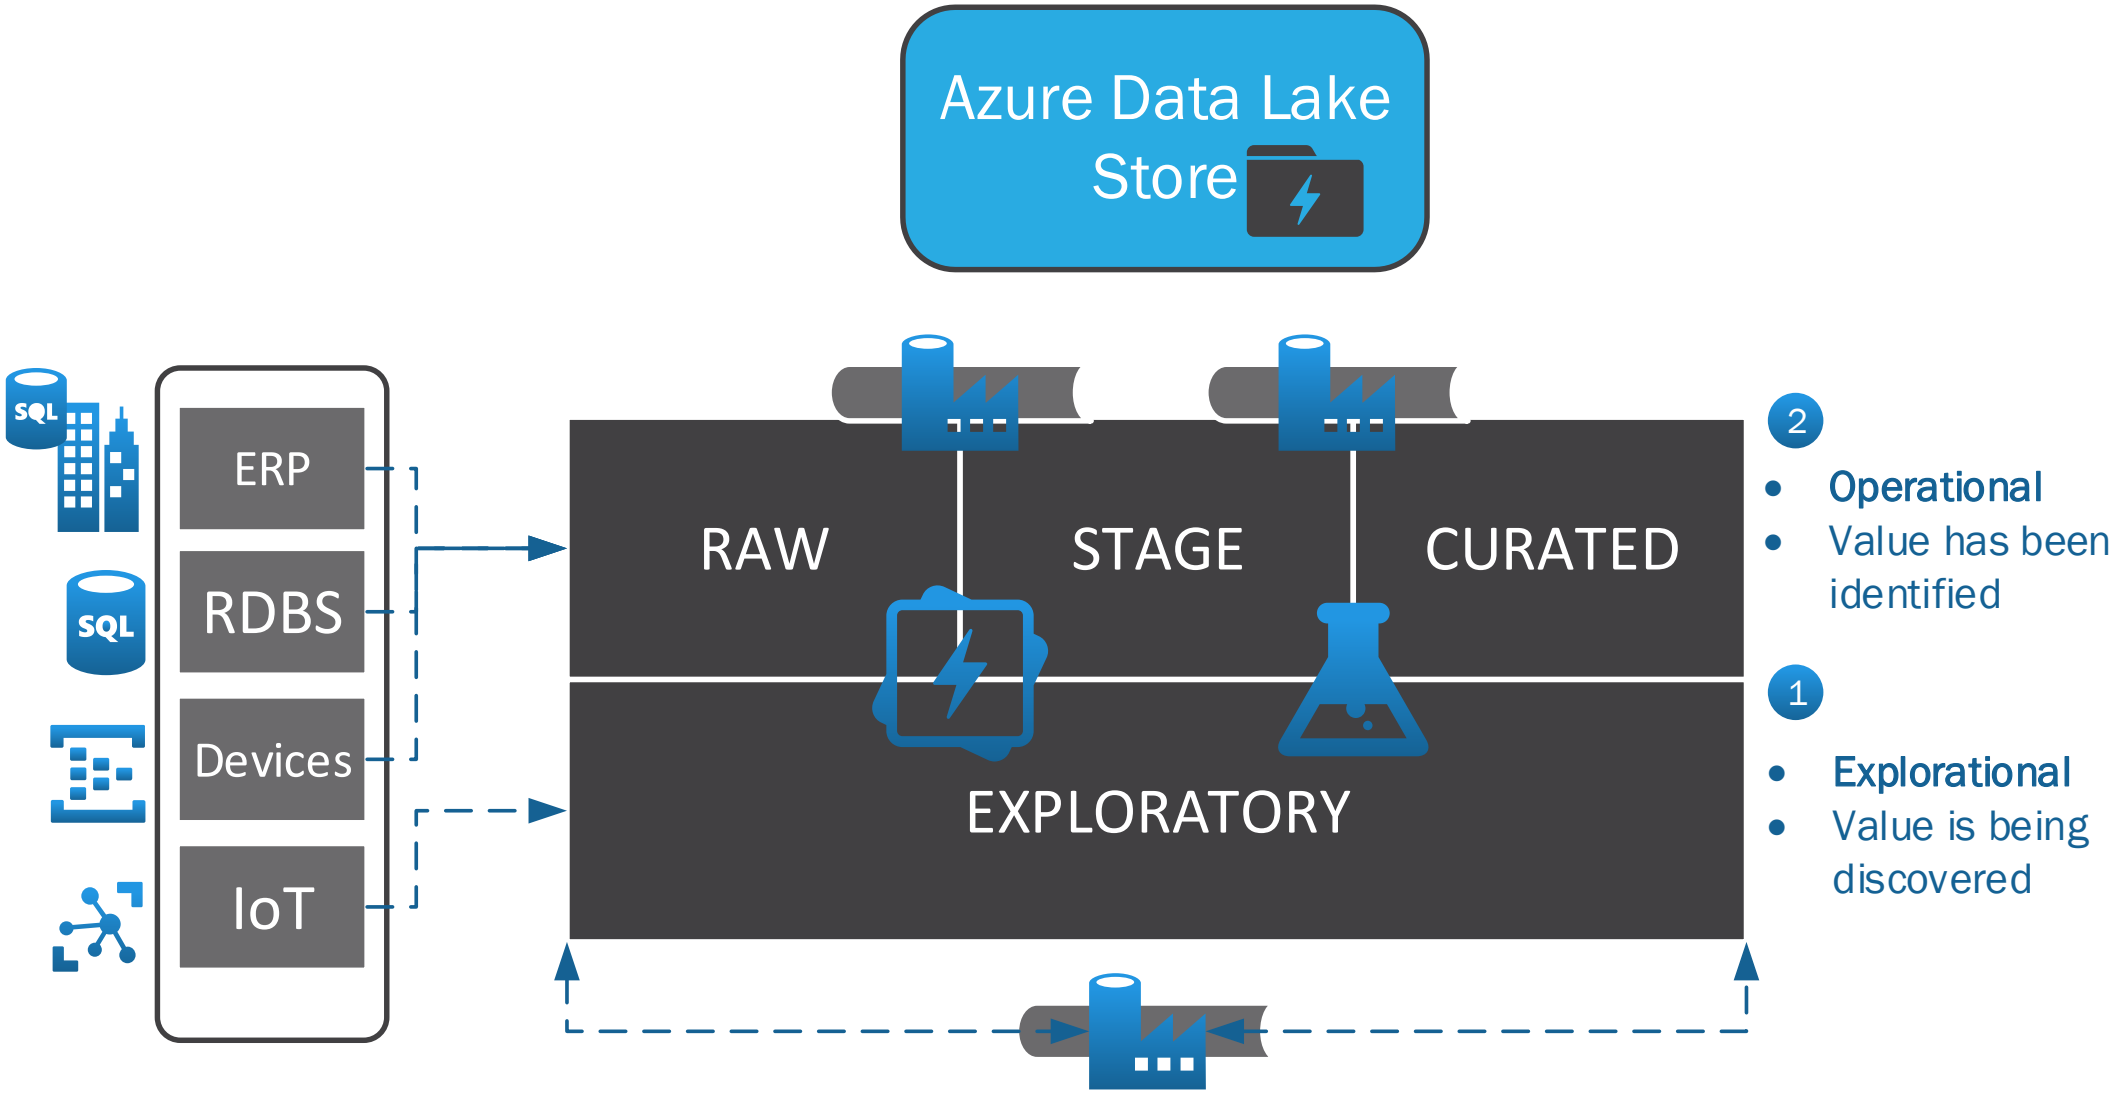

Data Lake Operations

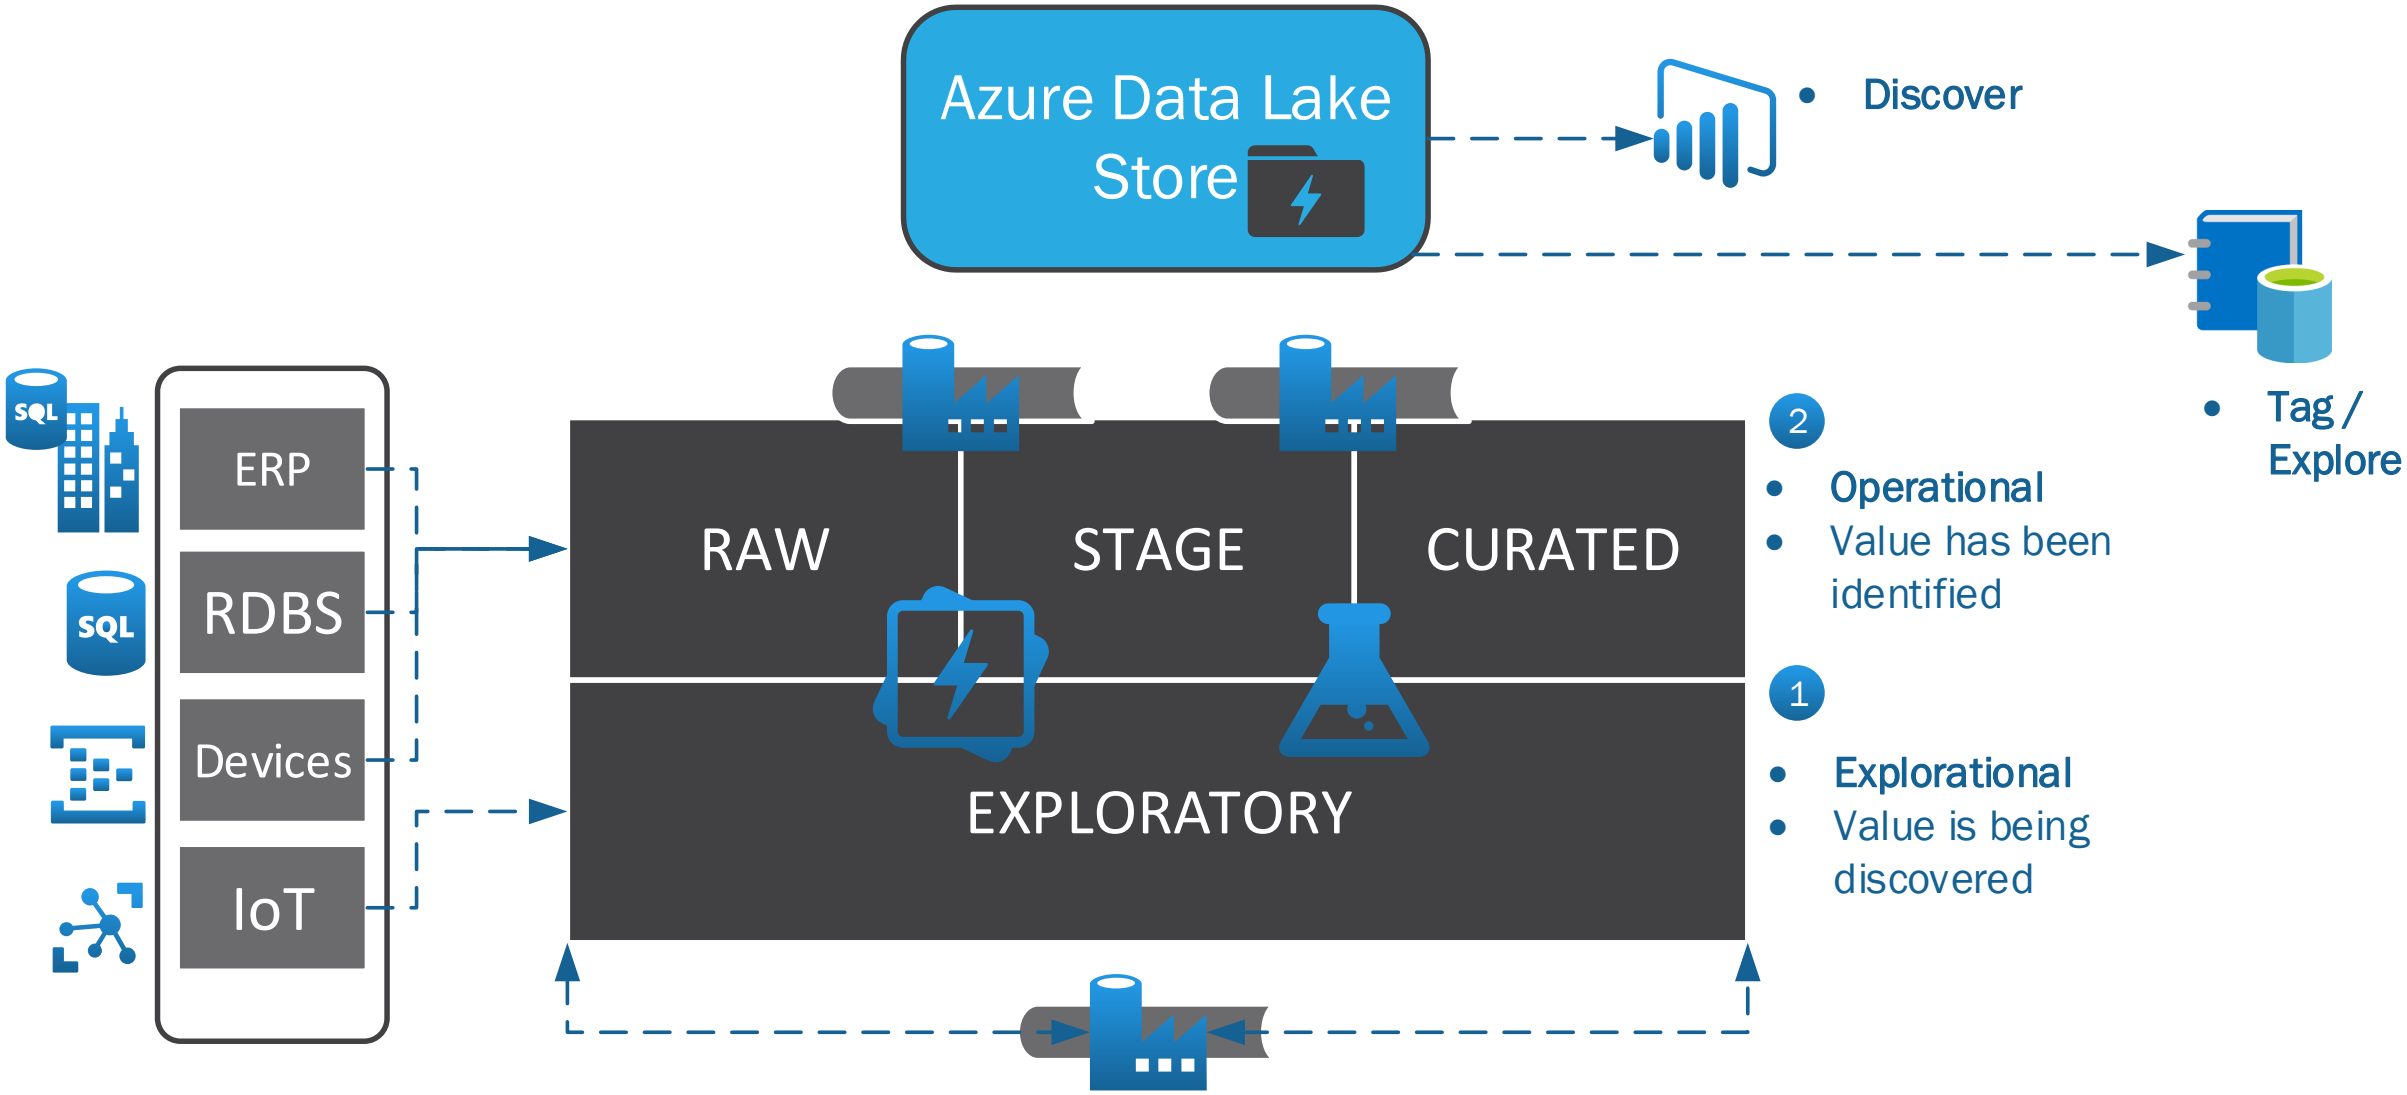

Data Lake Organization

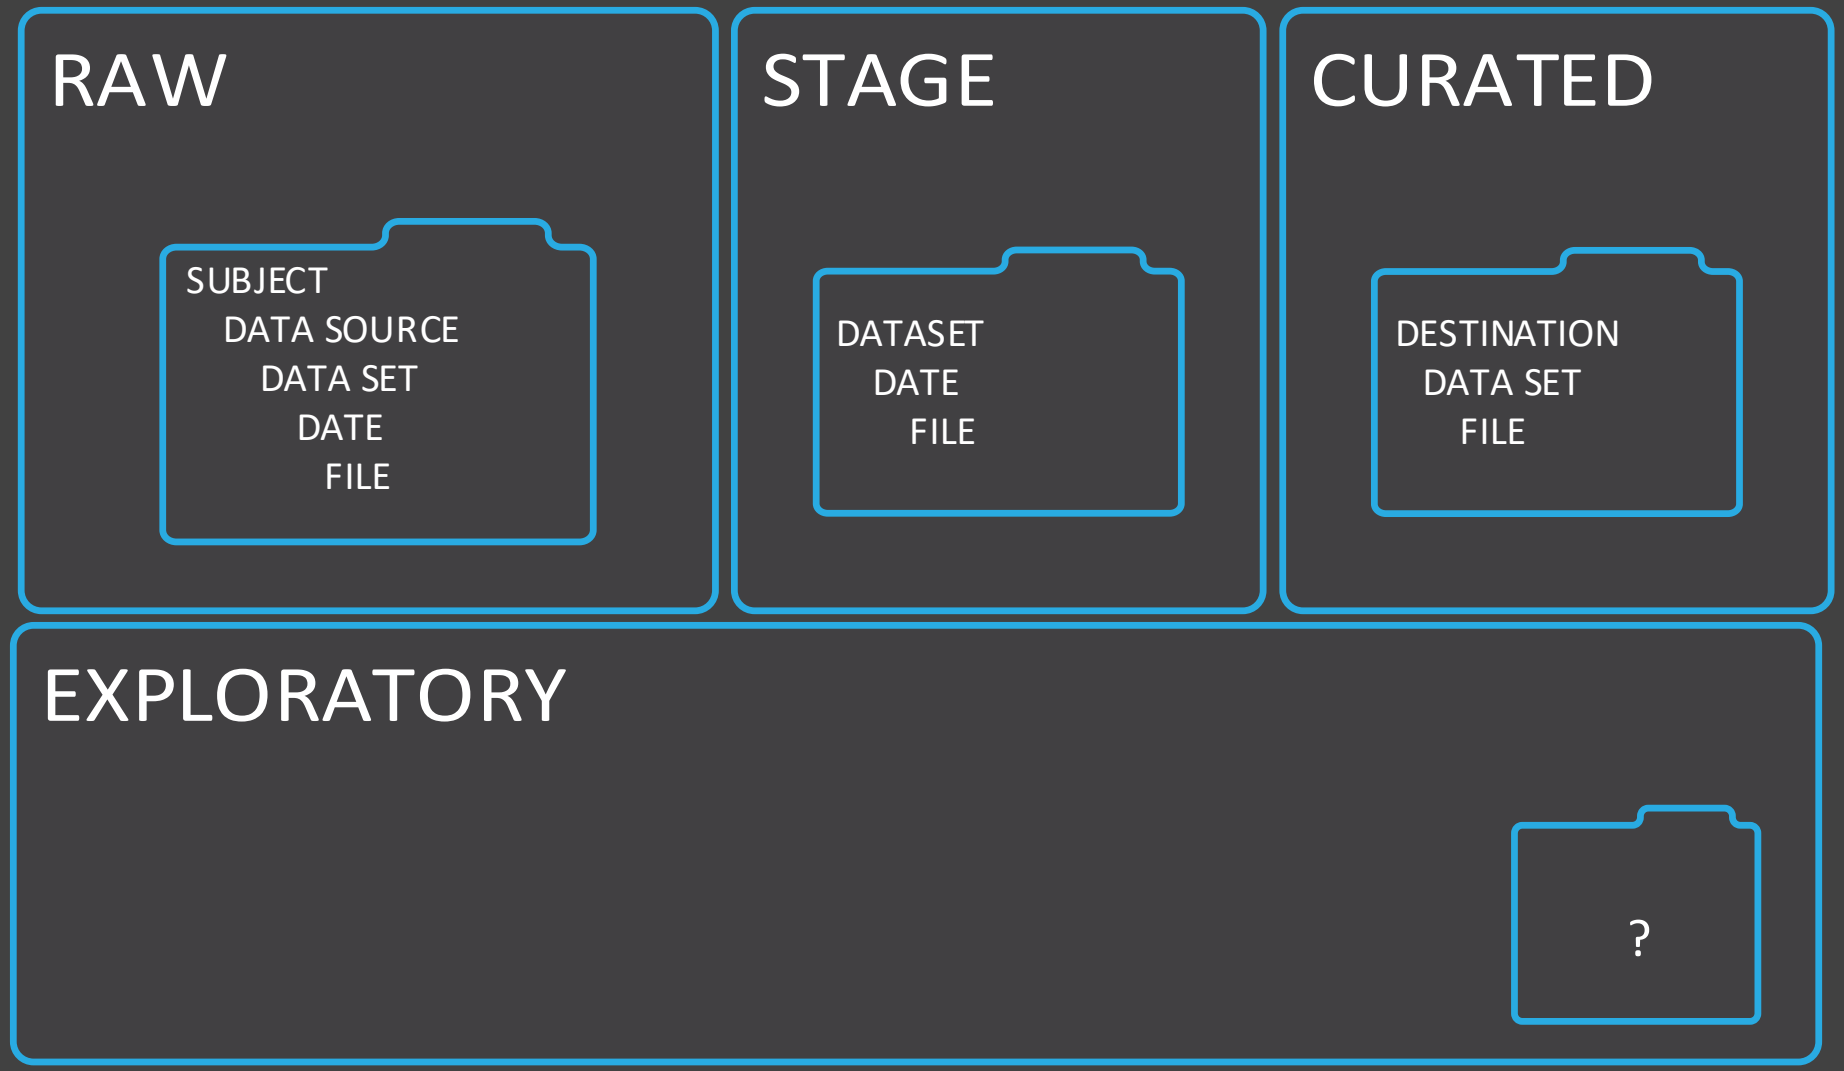

Data Lake Security

RAW (1)	STAGE (2)	CURATED (3)	EXPLORATION (0)
Data Experts/Engineers	Data Experts/Engineers	ETL and BI Engineers / SME's / Analysts	Data Scientist / Analysts

TOOLS

Data Lake Tagging

RAW (1)	STAGE (2)	CURATED (3)	EXPLORATION (0)
Data Experts/Engineers	Data Experts/Engineers	ETL and BI Engineers / SME's / Analysts	Data Scientist / Analysts
AUTOMATED		**SME**	**N/A**

TOOLS

Data Lake Processing

RAW (1)	STAGE (2)	CURATED (3)	EXPLORATION (0)
Data Experts/Engineers	Data Experts/Engineers	ETL and BI Engineers / SME's / Analysts	Data Scientist / Analysts
AUTOMATED		SME	N/A
INGESTION	**CLEANSING**	**DISTRIBUTION**	

TOOLS

Conceptual Architecture

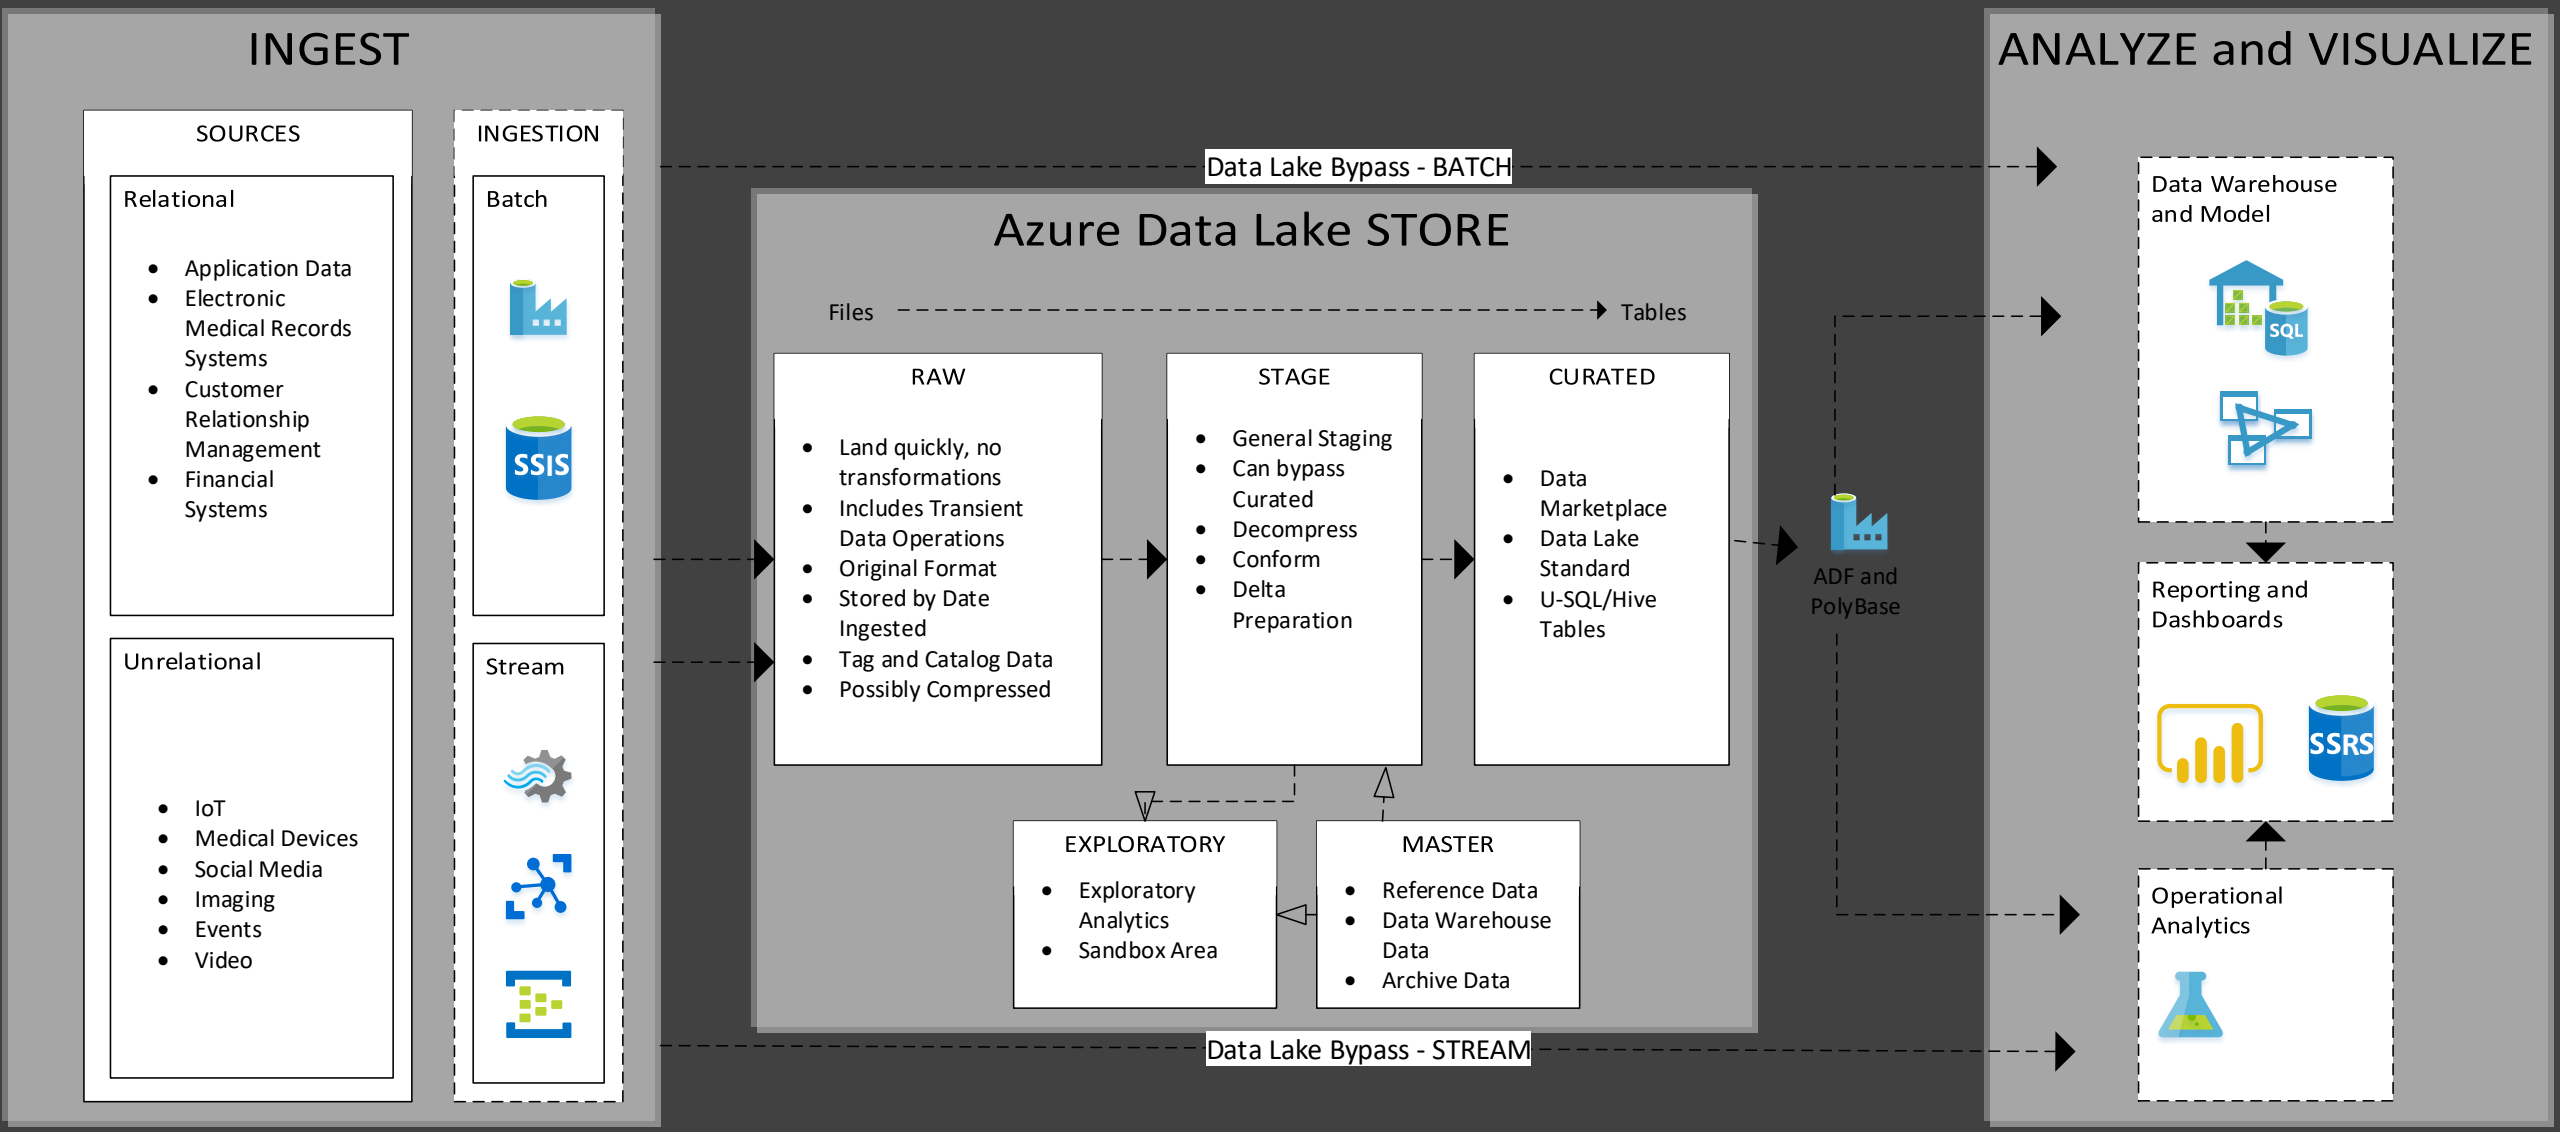

Data Lake Questions

1. Big Data Overview
2. Data Lake Concepts
3. **Azure Data Lake Store**
4. Azure Data Lake Analytics
5. U-SQL

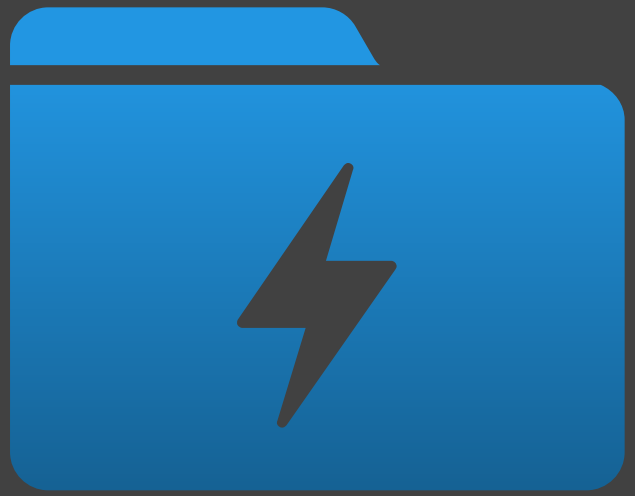

Azure Data Lake Store

- **HDFS** Data Store for housing data in it's Native **Raw** Format
  - Built on Apache YARN
- Process and store **Petabyte** size files
- Enterprise **Security** through Azure Active Directory
- No storage limit

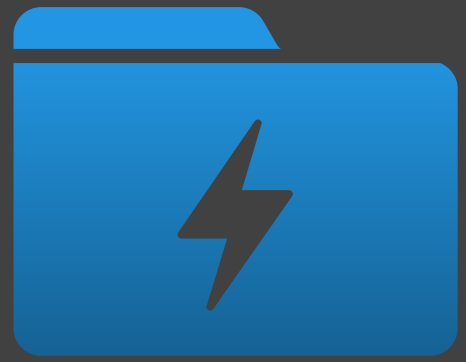

Data Lake Store

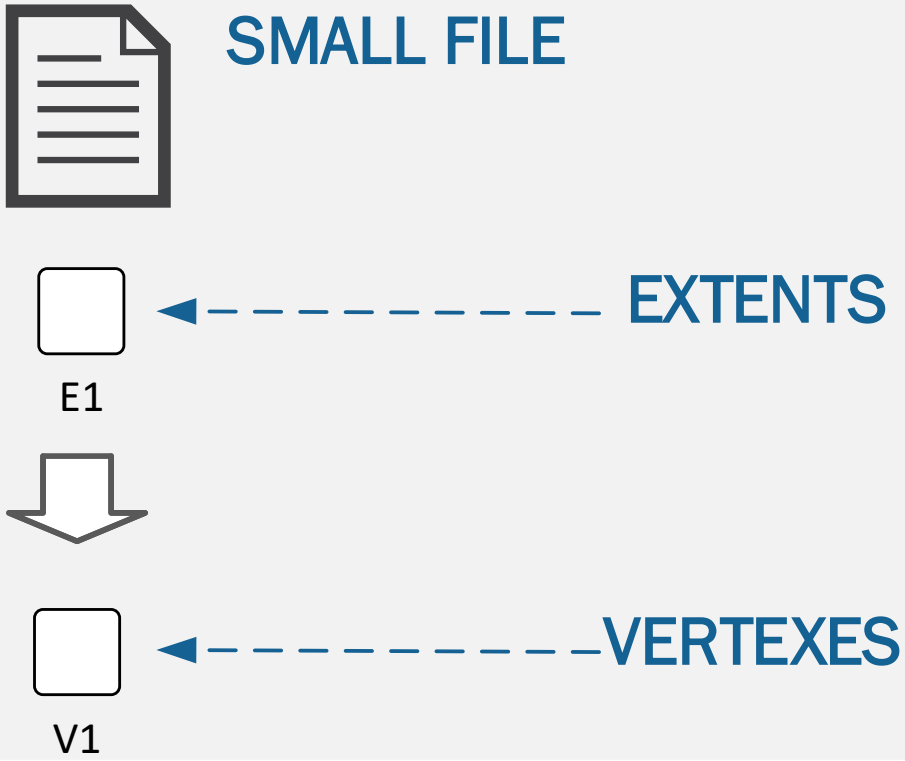

Data Lake Store

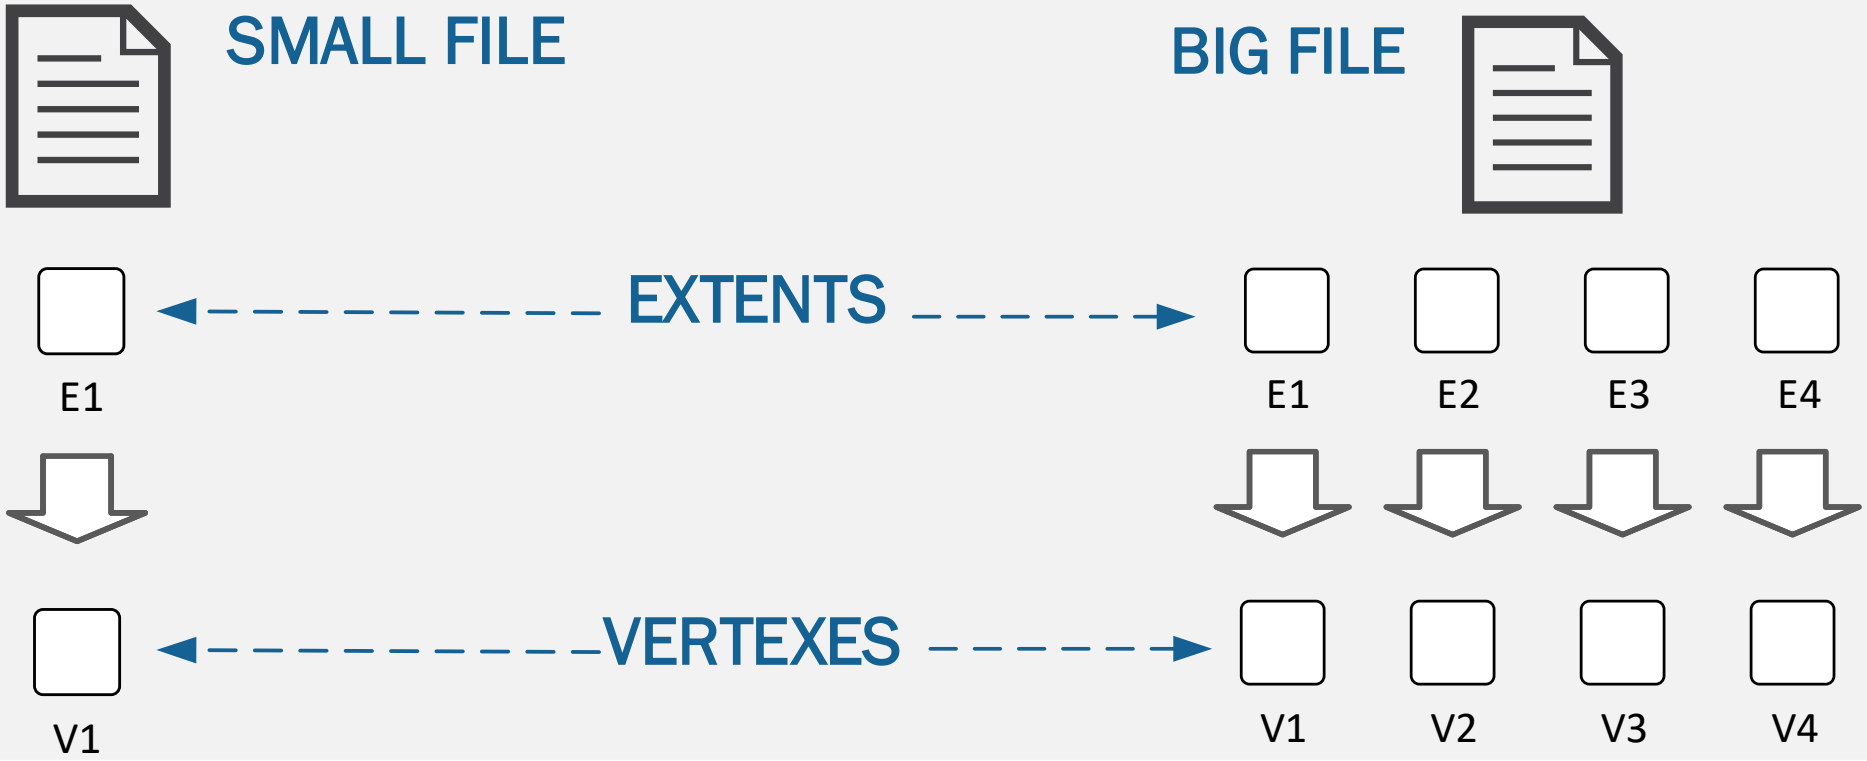

Data Lake Store

SMALL FILE

EXTENTS

E1

VERTEXES

V1

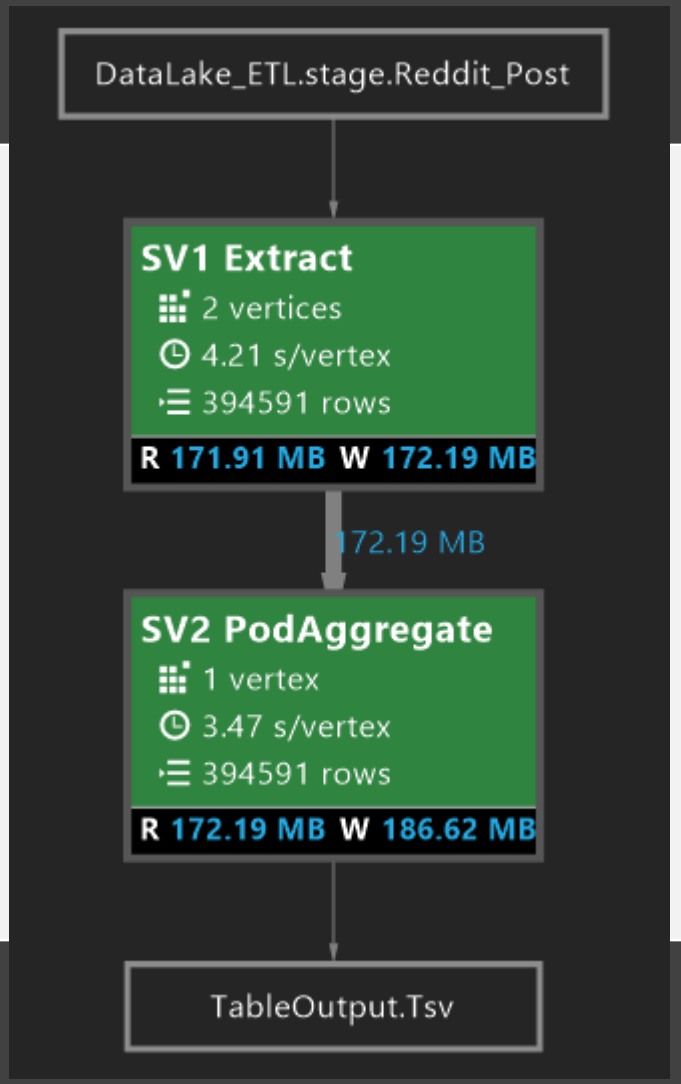

Data Lake Ingestion

- Visual Studio
- Azure Portal (Limit)
- SSIS Data Lake Destination and SSIS Data Lake Task
- Azure Data Factory
- Powershell
- ADLCopy
- Sqoop (HDInsight)
- Other Apache tools

1. Big Data Overview
2. Data Lake Concepts
3. Azure Data Lake Store
4. **Azure Data Lake Analytics**
5. U-SQL

Azure Data Lake Analytics

- **Big Data Queries** as a **Service**
- Analytics **Federation**
- Develop in **U-SQL**, **.NET**, **R**, and **Python**
- **Cognitive Services**
- **Scale** Instantly
- Pay **Per Job**

1. Big Data Overview
2. Data Lake Concepts
3. Azure Data Lake Store
4. Azure Data Lake Analytics
5. **U-SQL**

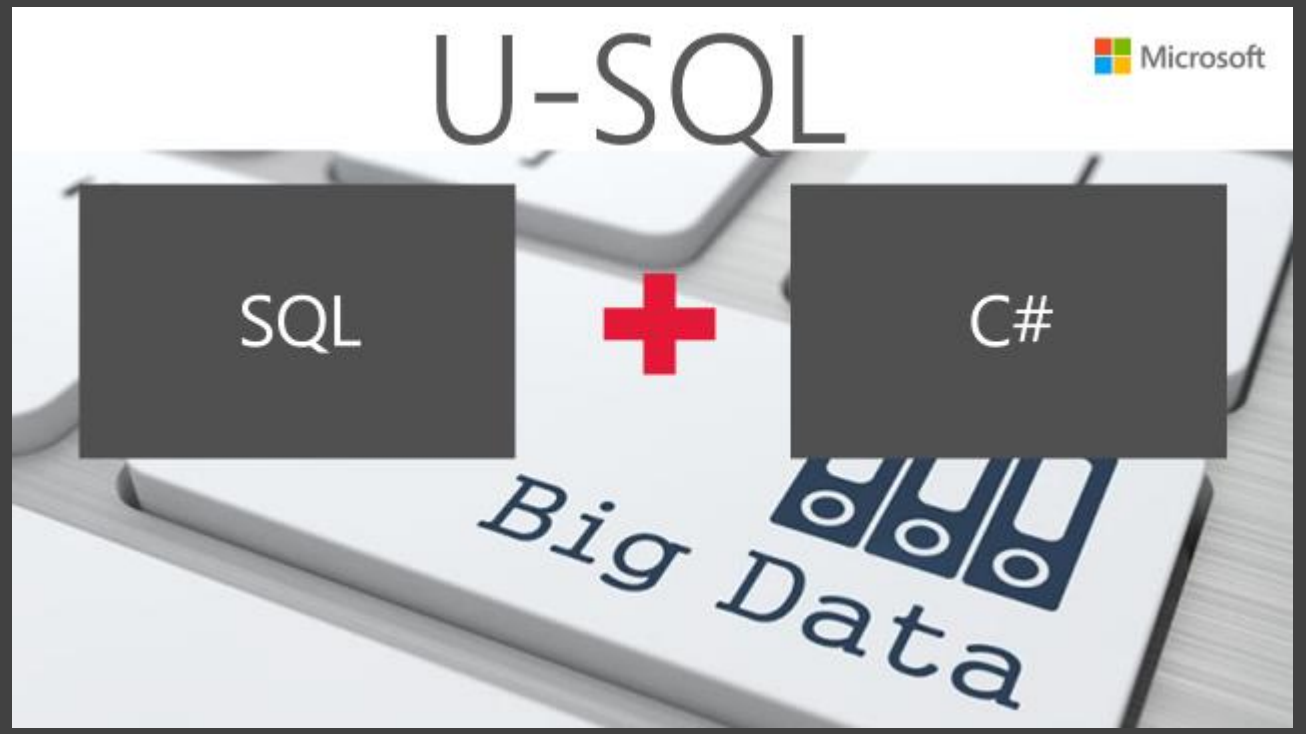

Intro to U-SQL

KEY FEATURES

- Combines **SQL** and **C#**
- **Patterned File Processing**
- Extensions: **Python**, **R**, **Cognitive**
- Query Data where it Lives (**Federated Querying**)
- **Partition** and **Distribution** of Data for **Massive Parallelism**
- Manage Structure and Shared Programming through **U-SQL Catalog**
- U-SQL Procedures

U-SQL : Extract Query

U-SQL

```
1 @MyExtract =
  EXTRACT
    Field1 string,
    Field2 int,
    Field 3 int?
  FROM "/datalake/01_RAW/{*}.csv"
  USING Extractors.Csv();
```

T-SQL

```
CREATE TABLE myTable
(
Field1 VARCHAR(100),
Field2 INT,
Field3 INT NOT NULL
);


INSERT INTO myTable
( Field1, Field2, Field3)
SELECT
  CAST(Field1 as varchar(100) as Field1,
  CAST(Field2 AS INT) as Field2,
  CONVERT(INT, Field3) as Field 3
FROM myTable
```

U-SQL : Extract Query

U-SQL

1
```
@MyExtract =
EXTRACT
  Field1 string,
  Field2 int,
  Field 3 int?
FROM "/datalake/01_RAW/{*}.csv"
USING Extractors.Csv();
```

2
```
@MyAgg =
SELECT
  Field1,
  MAX(Field2) A
FROM @MyExtract
GROUP BY Field1;
```

T-SQL

```
CREATE TABLE myTable
(
Field1 VARCHAR(100),
Field2 INT,
Field3 INT NOT NULL
);
```

```
INSERT INTO myTable
( Field1, Field2, Field3)
SELECT
  CAST(Field1 as varchar(100) as Field1,
  CAST(Field2 AS INT) as Field2,
  CONVERT(INT, Field3) as Field 3
FROM myTable
```

U-SQL : Extract Query

U-SQL

1
```
@MyExtract =
EXTRACT
  Field1 string,
  Field2 int,
  Field 3 int?
FROM "/datalake/01_RAW/{*}.csv"
USING Extractors.Csv();
```

2
```
@MyAgg =
SELECT
  Field1,
  MAX(Field2) AS Field2
FROM @MyExtract
GROUP BY Field1;
```

3
```
OUTPUT @MyAgg
TO     datalake/02_STAGE/MyOutput.csv"
USING Outputters.Csv();
```

T-SQL

```
CREATE TABLE myTable
(
Field1 VARCHAR(100),
Field2 INT,
Field3 INT NOT NULL
);
```

```
INSERT INTO myTable
( Field1, Field2, Field3)
SELECT
  CAST(Field1 as varchar(100) as Field1,
  CAST(Field2 AS INT) as Field2,
  CONVERT(INT, Field3) as Field 3
FROM myTable
```

U-SQL : Extractors and Outputters

CURRENT EXTRACTORS

- Csv()
- Tsv()
- Txt()

CURRENT OUTPUTTERS

- Csv()
- Tsv()
- Txt()

U-SQL : Extractor and Outputter Parameters

```
EXTRACT
…
FROM “/datalake/01_RAW/{*}.CSV
USING Extractors.Csv(silent : true , delimiter : “,”)
```

(); PARAMETERS

- Delimiter
- Encoding
- escapeCharecter
- nullEscape
- Quoting
- rowDelimiter
- Silent
- skipFirstNRows
- charFormat

U-SQL : Extractors and Outputters

CURRENT EXTRACTORS

- Csv()
- Tsv()
- Txt()

CURRENT OUTPUTTERS

- Csv()
- Tsv()
- Txt()

CUSTOM EXTRACTORS and OUTPUTTERS

- FlexExtractor()
- XML()
- JSON()
- Avro()

U-SQL : Virtual Columns

```
DECLARE        =“/datalake/01_stage/2017/06/{FileName}.csv

@MyExtract =
EXTRACT
  Field1 string
  Field2 int,
  Field 3 int?,
  FileName
FROM @IN
USING Extractors.Csv()
WHERE FileName == “MyRedditFile_20170602”;
```

File Names :

MyRedditFile_20170601.csv
MyRedditFile_20170602.csv
etc…

U-SQL: Job Execution

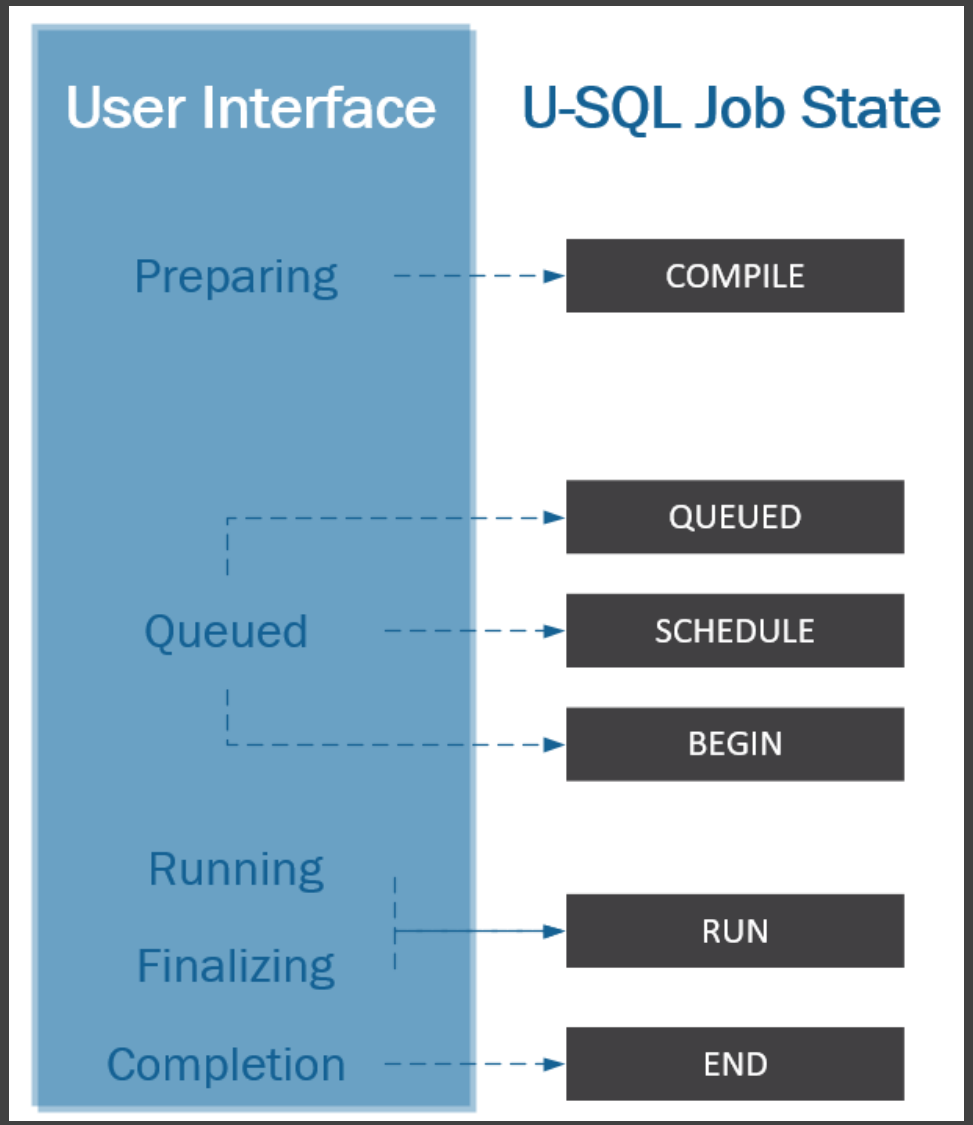

U-SQL: Job Execution

Preparing	Queued	Executing	Finalizing
29 seconds	7 seconds	20 seconds	11 seconds

Job Result	Succeeded
Total Duration	1.1 minutes
Submit Time	6/21/2017 10:55:04 PM
Start Time	6/21/2017 10:55:40 PM
End Time	6/21/2017 10:56:11 PM
Compilation	29 seconds
Queued	7 seconds
Running	31 seconds
Account	talavantanalytics
Author	sean.forgatch@talavant.com
Priority	1000
Parallelism	5
Bytes Left	60,978
Bytes Read	360,810,773
Bytes Written	376,235,827
Total Vertices	3
Completed	3
Running	0
Failed	0

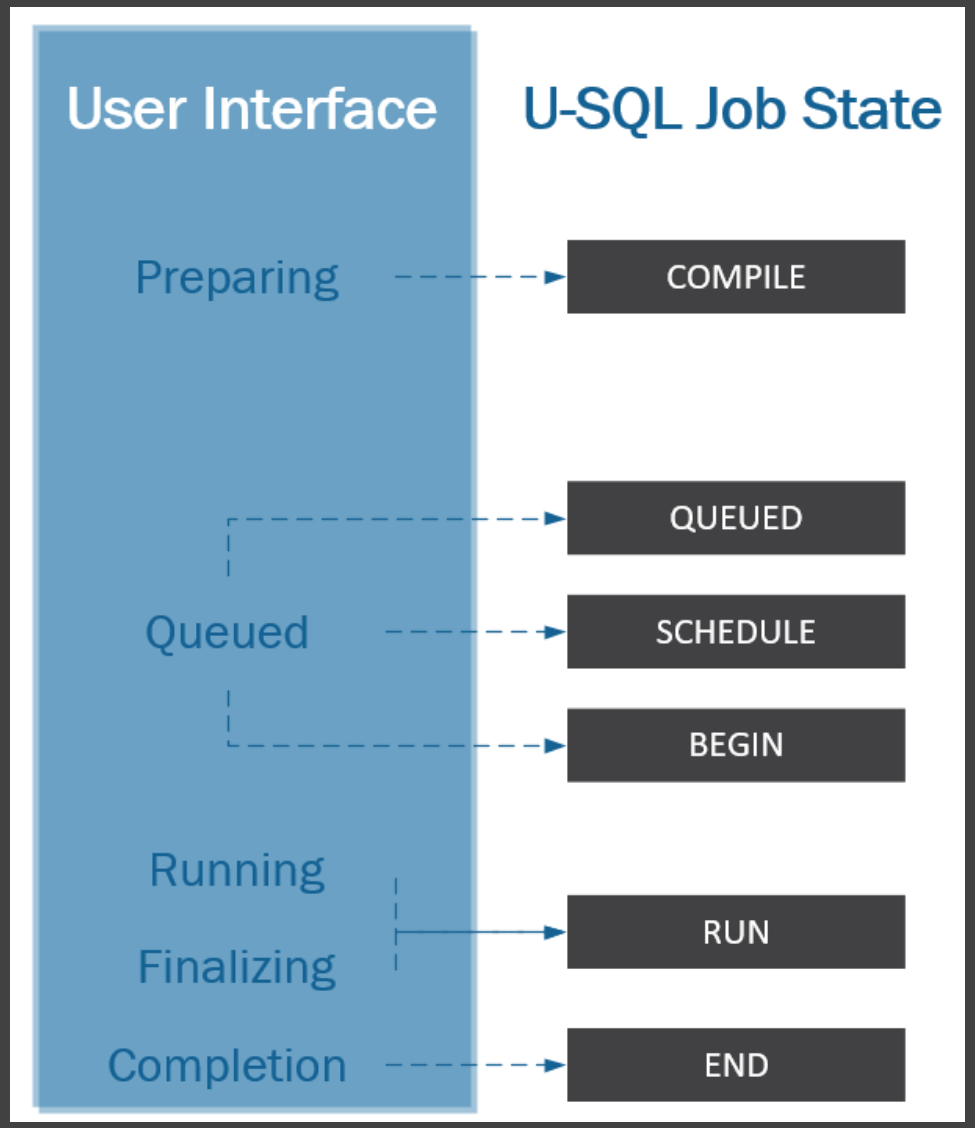

TALAVANT

U-SQL : Tables

GUIDELINES

1. Must Have **Clustered Index**
2. Utilize When **Improving Performance** with Distribution/Partitioning
3. You have **Multiple Large Files**
4. Don’t Use when:
   - No Filtering, Joining, Grouping

U-SQL : Tables

```
DROP TABLE IF EXISTS <adla>.<database>.<schema>.tableName;
CREATE TABLE <adla>.<database>.<schema>.tableName
(
Field1 int,
Field2 string,
Field3 int?

INDEX idx_1 CLUSTERED(Field1)
DISTRIBUTED BY HASH(Field2)
);
```

U-SQL : Tables

```
DROP TABLE IF EXISTS
<adla>.<database>.<schema>.tableName;
CREATE TABLE <adla>.<database>.<schema>.tableName
(
Field1 int,
Field2 string,
Field3 int?

INDEX idx_1 CLUSTERED(Field1)
DISTRIBUTED BY HASH(Field2)
)
```

AS SELECT …

U-SQL : Tables

```
DROP TABLE IF EXISTS
<adla>.<database>.<schema>.tableName;
CREATE TABLE <adla>.<database>.<schema>.tableName
(
Field1 int,
Field2 string,
Field3 int?

INDEX idx_1 CLUSTERED(Field1)
DISTRIBUTED BY HASH(Field2)
)
```

AS EXTRACT …

U-SQL : Tables

```
DROP TABLE IF EXISTS
<adla>.<database>.<schema>.tableName;
CREATE TABLE <adla>.<database>.<schema>.tableName
(
Field1 int,
Field2 string,
Field3 int?

INDEX idx_1 CLUSTERED(Field1)
DISTRIBUTED BY HASH(Field2)
)
```

AS TVF …

U-SQL : Tables

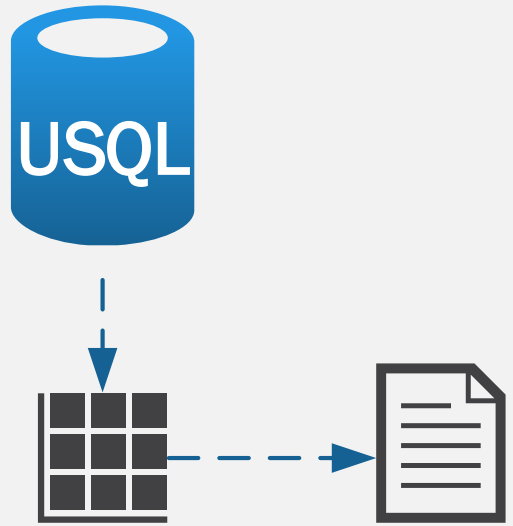

METADATA and DATA

MANAGED

- Own Their Data
- No Heaps
- INSERT only

U-SQL : Tables

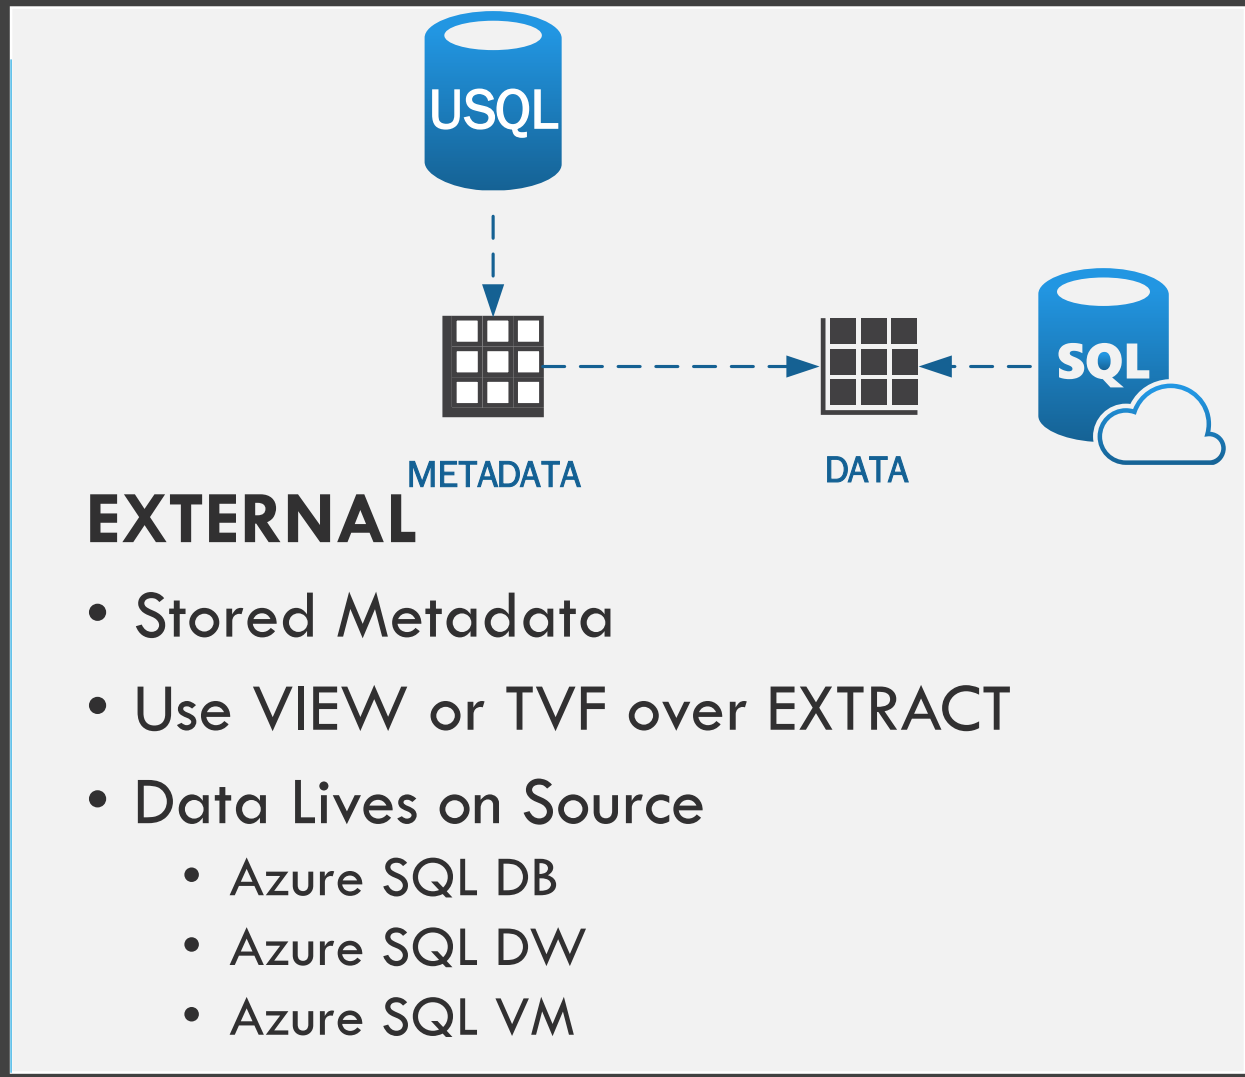
USQL
METADATA
DATA
SQL
EXTERNAL
• Stored Metadata
• Use VIEW or TVF over EXTRACT
• Data Lives on Source
  • Azure SQL DB
  • Azure SQL DW
  • Azure SQL VM

U-SQL : Data Federation

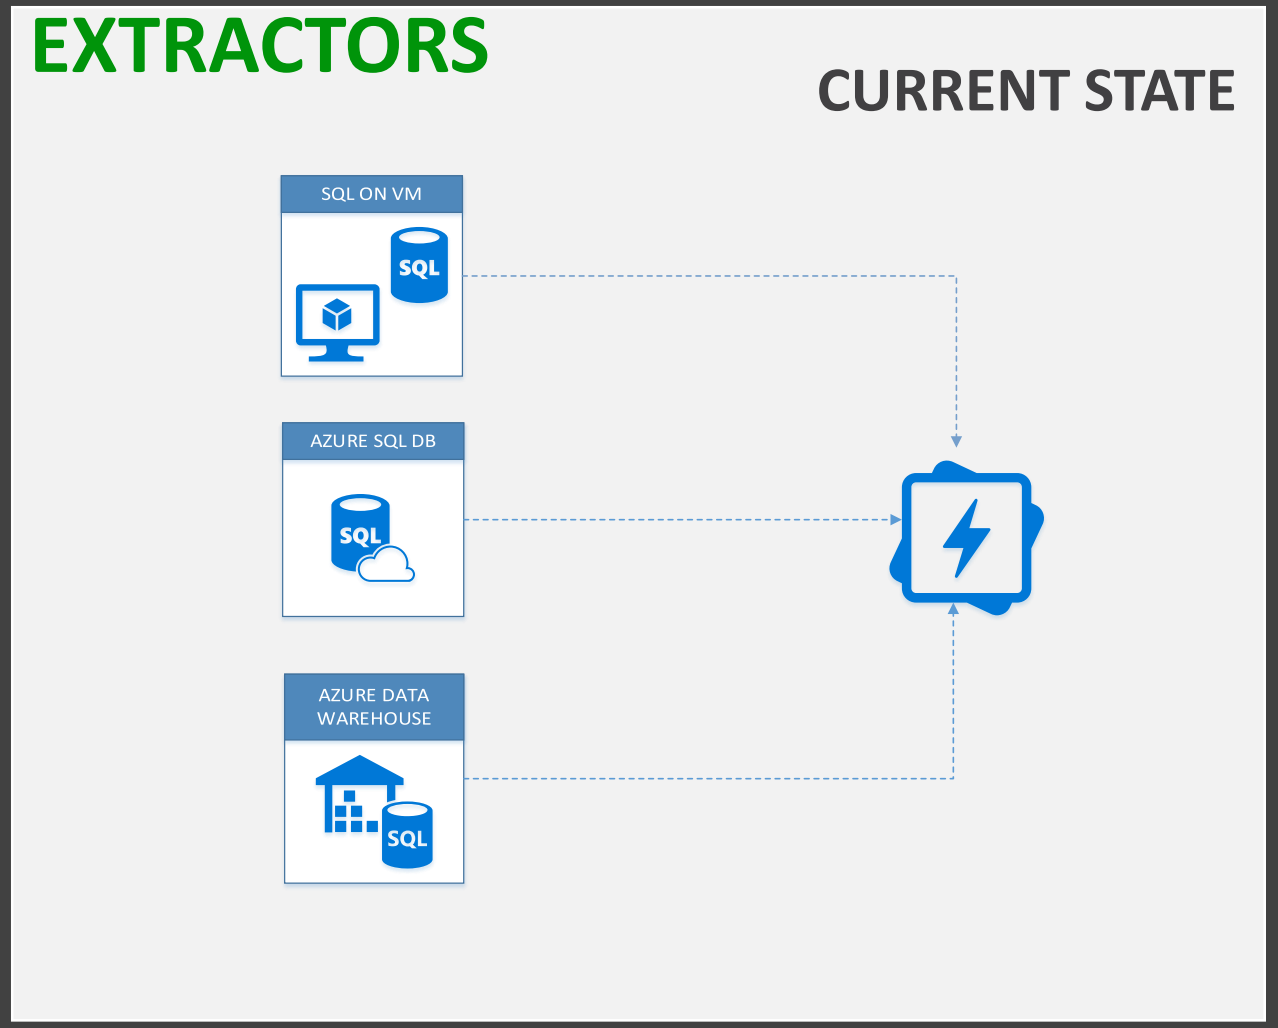

U-SQL : Operators

COMPARISON OPERATORS	LOGICAL OPERATORS
IS NULL	AND
==	BETWEEN
>	IN, NOT IN
>=	LIKE, NOT LIKE
!=	NOT
	OR

U-SQL : Functions

REPORTING FUNCTIONS

- COUNT
- SUM
- MIN
- MAX
- AVG
- STDEV
- VAR

ANALYTIC FUNCTIONS

- CUME_DIST
- PERCENT_RANK
- PERCENTILE_CONT
- PERCENTILE_DISC

RANKING FUNCTIONS

- RANK
- DENSE_RANK
- NTILE
- ROW_NUMBER

U-SQL : C# Functions

MATH METHODS

- Abs
- BigMul
- Floor
- Max/Min
- Round
- Sqrt
- ..plus many more!

STRING METHODS

- Compare
- Concat
- Contains
- Equals
- Replace
- Split
- ToUpper
- Trim
- ..plus many more!

Advice

- **Identify Value of Data Lake Approach**
- **Data Lake: Invest Time and Strategy into Data Lake Design**
- **U-SQL: Utilize U-SQL Constructs before C#**
- **U-SQL: Understand and Control Data through Partitioning**

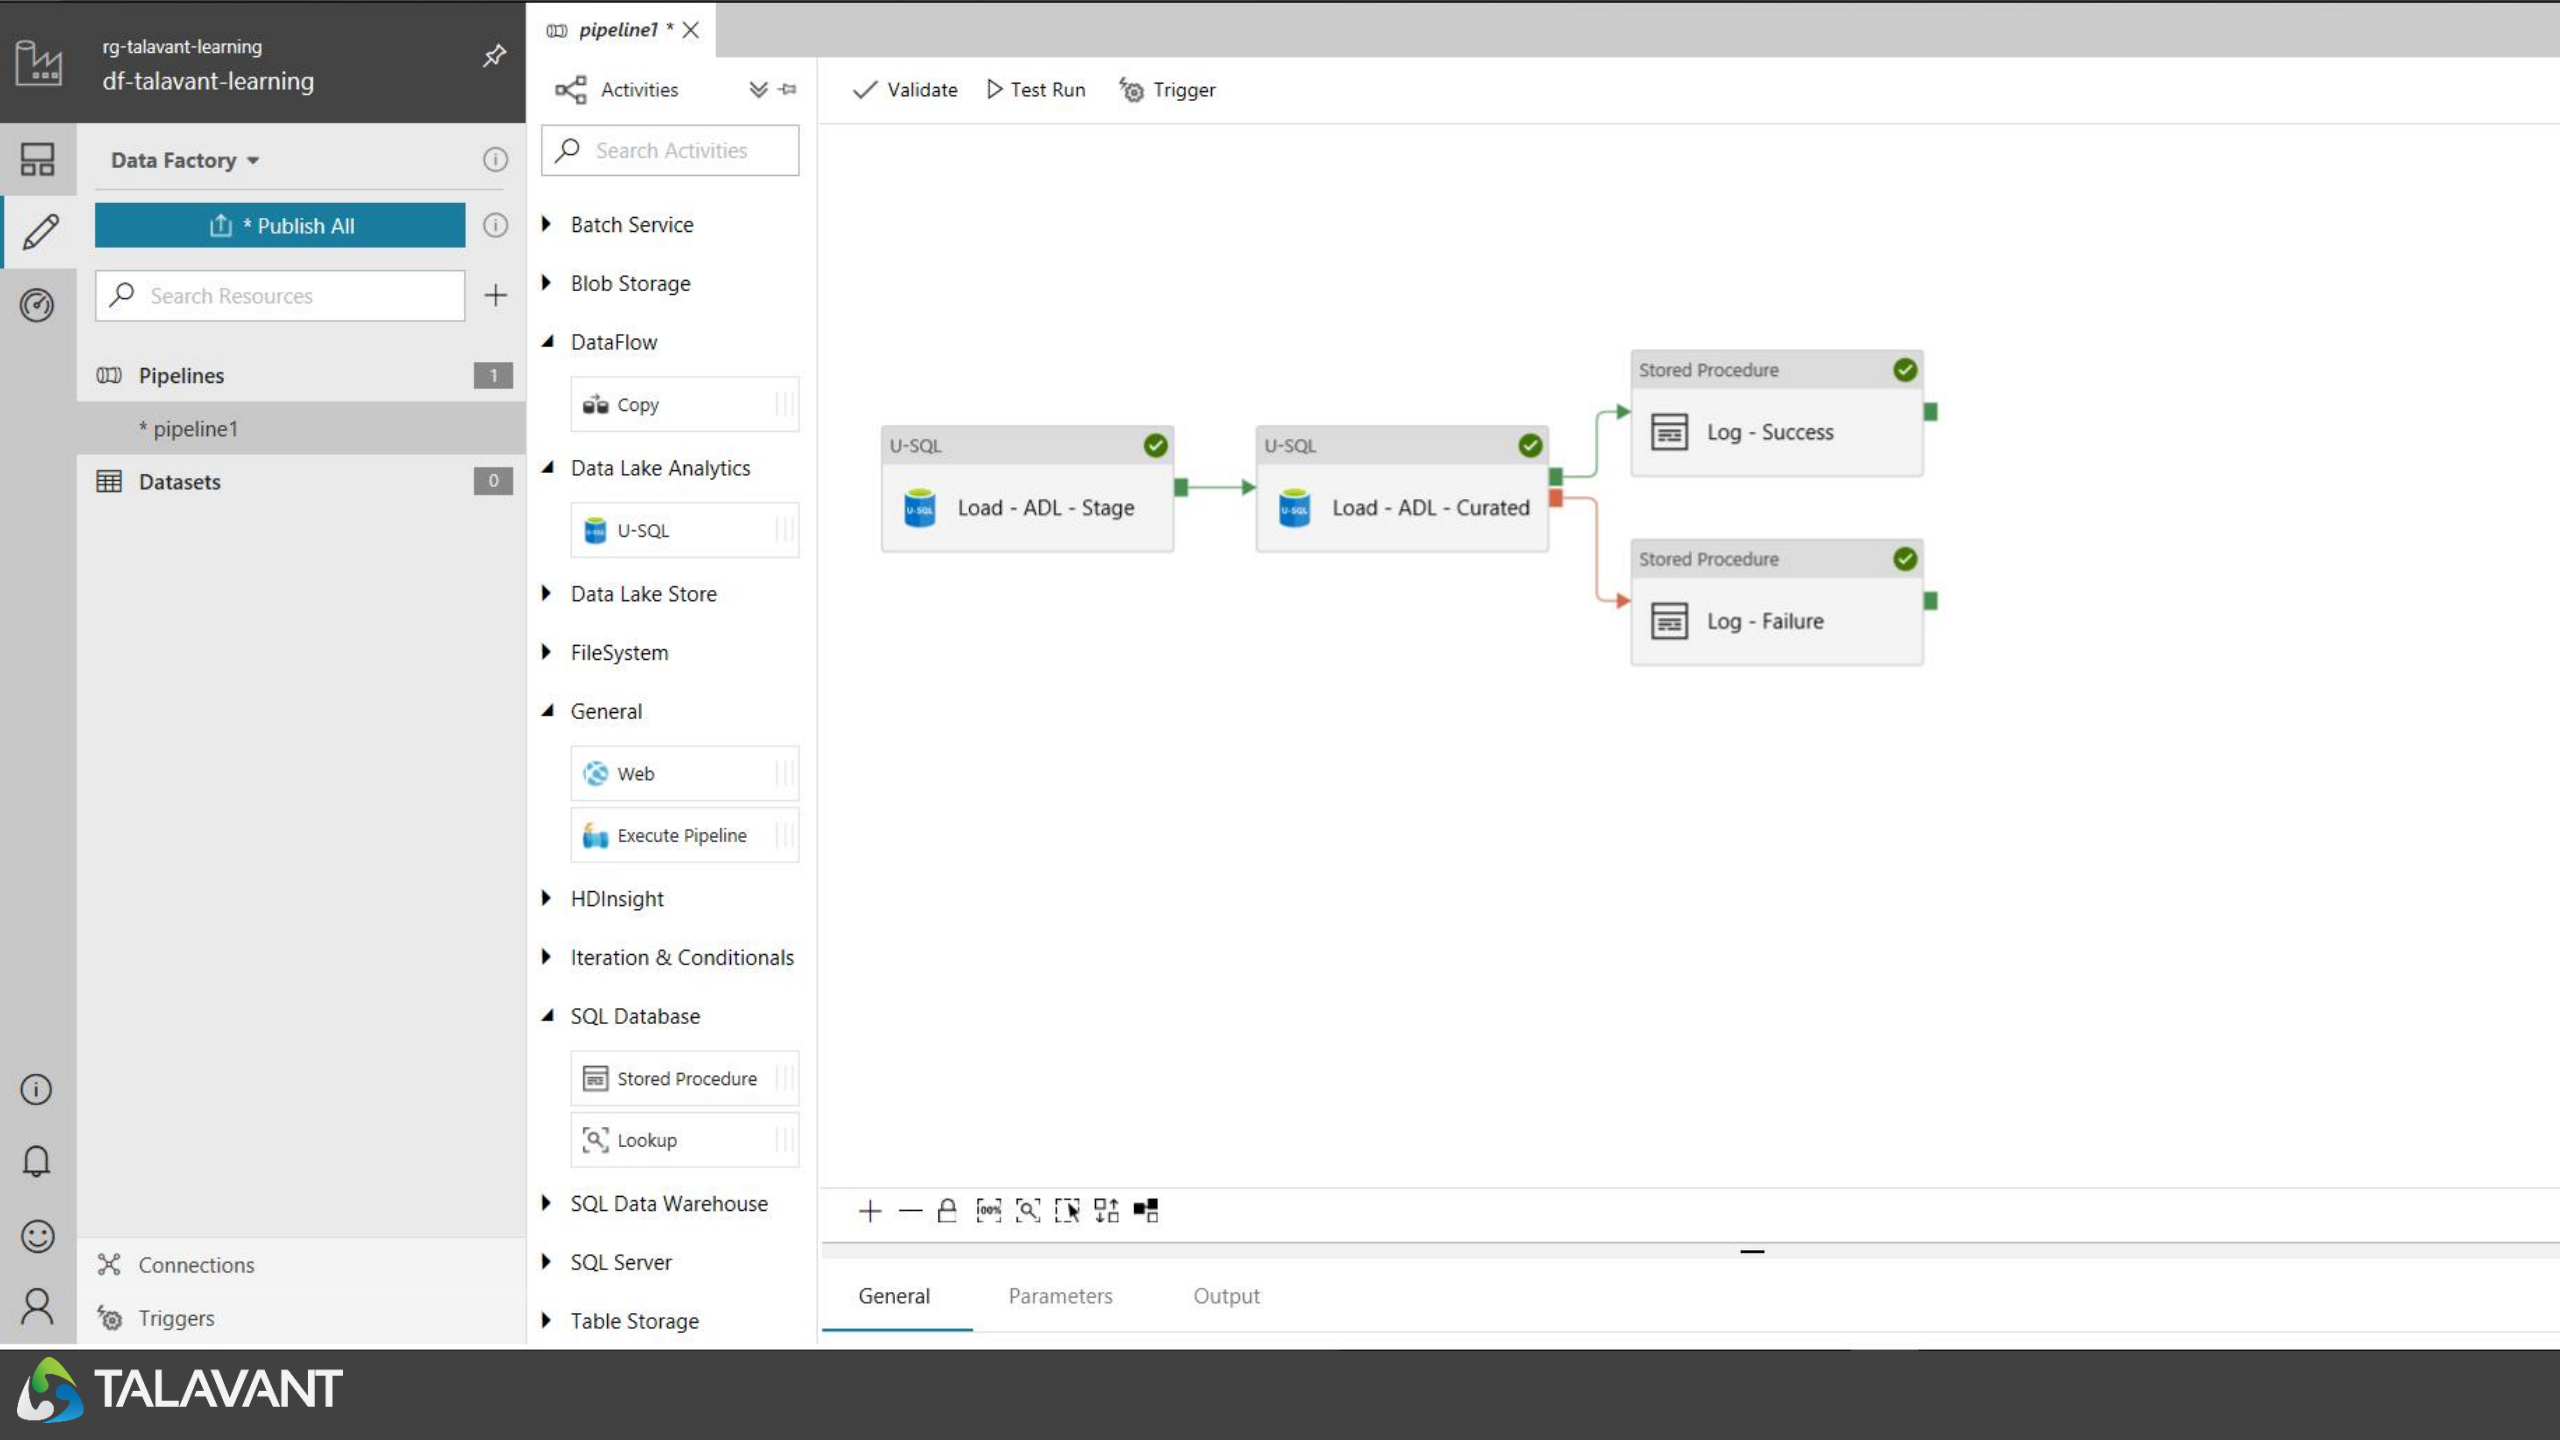
rg-talavant-learning
df-talavant-learning
pipeline1 *
Activities
Validate
Test Run
Trigger
Data Factory
* Publish All
Search Resources
Search Activities
Pipelines 1
* pipeline1
Datasets 0
Batch Service
Blob Storage
DataFlow
Copy
Data Lake Analytics
U-SQL
Data Lake Store
FileSystem
General
Web
Execute Pipeline
HDInsight
Iteration & Conditionals
SQL Database
Stored Procedure
Lookup
SQL Data Warehouse
SQL Server
Table Storage
U-SQL
Load - ADL - Stage
U-SQL
Load - ADL - Curated
Stored Procedure
Log - Success
Stored Procedure
Log - Failure
Connections
Triggers
General
Parameters
Output

Learn U-SQL !

- **Michael Rys** – LinkedIn Slide Share's
- **GitHub** – U-SQL Repository
- SQL Server Central – **Stairway to U-SQL**
- **Azure** – Built in Example

Let's Connect!

https://www.linkedin.com/in/seanforgatch/

Sean.Forgatch@Talavant.com

- @4gatchSQL

TALAVANT

PASS